

Architecture of Massively Scalable Applications, Spring 2026

Milestone 2: Freelance Marketplace Platform
Polyglot Persistence, Authentication, Caching & Design Patterns
Deadline is Sunday 07/06/2026 at 11:59 PM

Contents

1 Overview	4
2 Git Workflow — Feature Development	5
3 Design Patterns	7
3.1 The 7 Required Patterns	7
3.2 Strategy Pattern — Payout Reversal Logic (S5-F12)	7
3.3 Observer Pattern — Event Logging.....	8
3.4 Chain of Responsibility — JWT Filter Chain	9
3.5 Builder Pattern — Dashboard DTOs.....	10
3.6 Singleton Pattern — JwtConfigurationManager.....	11
3.7 Factory Pattern — Event Creation	12
3.8 Adapter Pattern — NoSQL Result → DTO Conversion	13
3.9 Grading Summary	13
4 M1 Modifications Required	14
4.1 User Entity — Password Hashing.....	14
4.2 User Entity — Role Values (Additive).....	14
4.3 Existing M1 Endpoints — JWT Authentication.....	15
4.4 Existing M1 Read Endpoints — Redis Caching (Explicit Enumeration)	16
4.5 Design Pattern Retrofits to M1	19
4.6 Payout.transactionDetails — Additive platformFee Key	21
4.7 Summary Checklist	21
5 Authentication & Authorization	21
5.1 Overview.....	21
5.2 JWT Token	22
5.3 Roles (Additive — Preserve M1 Values)	22
5.4 Security Configuration (Per Service)	22
5.5 Password Handling	23
6 Database Architecture	23
6.1 Six Databases (Versions Match Task 2 Exactly).....	23

6.2	Memory Limitations (Required).....	23
6.3	Dependency Model	23
6.4	Docker Compose Additions	24
6.5	Reference application.yml Fragments	25
7	New Entity Models	26
7.1	MongoDB Document Entities (All Services)	26
7.1.1	Common MongoEvent Interface.....	27
7.1.2	AuthEvent (User Service).....	27
7.1.3	JobEvent (Job Service)	27
7.1.4	ProposalEvent (Proposal Service)	27
7.1.5	ContractEvent (Contract Service)	28
7.1.6	PayoutAuditEvent (Wallet Service).....	28
7.2	Elasticsearch Document (Job Service).....	29
7.2.1	JobSearchDocument	29
7.3	Neo4j Entities (Proposal Service).....	29
7.3.1	FreelancerNode.....	29
7.3.2	JobNode.....	29
7.3.3	PROPOSED_TO Relationship	29
7.4	Cassandra Entity (Contract Service)	30
7.4.1	ContractMilestoneEvent	30
8	Caching Strategy	30
8.1	TTL Guidelines	30
8.2	Invalidation	30
9	Cross-Cutting Requirements	30
9.1	CC-1 — JWT on All Endpoints.....	30
9.2	CC-2 — Role Management.....	31
9.3	CC-3 — Redis Caching on M1 Endpoints	32
9.4	CC-4 — Design Pattern Implementation	32
9.5	CC-5 — Docker Compose with 6 Databases.....	32
9.6	CC-6 — Application Configuration (application.yml).....	33
10	Features	33
10.1	User Service Features (port 8081)	33
10.1.1	[S1-F10] Register User.....	33
10.1.2	[S1-F11] Login	34
10.1.3	[S1-F12] Get User Activity Feed	35
10.2	Job Service Features (port 8082)	36
10.2.1	[S2-F10] Full-Text Job Search.....	36
10.2.2	[S2-F11] Index Job for Search	37
10.2.3	[S2-F12] Get Job Market Dashboard.....	38

10.3	Proposal Service Features (port 8083)	39
10.3.1	[S3-F10] Get Proposal Analytics Dashboard	39
10.3.2	[S3-F11] Record Freelancer-Job Interaction	39
10.3.3	[S3-F12] Get Recommended Jobs for Freelancer	40
10.4	Contract Service Features (port 8084)	41
10.4.1	[S4-F10] Get Contract Analytics Dashboard	41
10.4.2	[S4-F11] Record Contract Milestone Event	42
10.4.3	[S4-F12] Get Contract Milestone Timeline	43
10.5	Wallet Service Features (port 8085)	44
10.5.1	[S5-F10] Get Platform Fee Analytics by Job Category	44
10.5.2	[S5-F11] Get Payout Method Breakdown	45
10.5.3	[S5-F12] Process Milestone-Based Payout Reversal	46

1 Overview

This milestone is worth **15%** of your final grade. You will extend the Freelance Marketplace Platform you built in Milestone 1 by adding **authentication, five NoSQL databases, caching, and design patterns**. Your existing 5 services, PostgreSQL database, and 45 features remain intact but require modifications (see Section 4) — Milestone 2 adds new capabilities on top of them.

What you are adding (four pillars):

- a) **Design Patterns** — 7 design patterns from Lab 7 applied at well-defined locations across your codebase (Strategy, Observer, Chain of Responsibility, Builder, Singleton, Factory, Adapter). See Section 3.
- b) **Authentication & Authorization** — JWT-based authentication with BCrypt password hashing and role-based access control. The User Service becomes the auth provider; all other services validate tokens independently.
- c) **NoSQL Databases** — MongoDB, Redis, Elasticsearch, Neo4j, and Cassandra. Each database serves a specific architectural role.
- d) **Caching** — Redis-based caching on all read-heavy endpoints, including your existing Milestone 1 endpoints.

You will also modify your existing M1 codebase (see Section 4) to integrate password hashing, JWT protection, caching, and design pattern retrofits.

Technologies (new in M2): Spring Security, JWT (JJWT), BCrypt, MongoDB, Redis, Elasticsearch, Neo4j, Cassandra, Spring Data for each database.

Spring Boot 4.0.3 on JDK 25 (Docker base image: eclipse-temurin:25.0.2_10-jdk). Jackson dual dependency: Jackson 3.x (tools.jackson.*) for Spring Boot, Jackson 2.x (com.fasterxml.*) for Hibernate 7.2's JSONB FormatMapper only.

Services (unchanged from M1):

- a) **User Service** (port 8081) — Now also the authentication provider
- b) **Job Service** (port 8082) — Now with full-text search via Elasticsearch
- c) **Proposal Service** (port 8083) — Now with recommendation graph via Neo4j
- d) **Contract Service** (port 8084) — Now with time-series milestone tracking via Cassandra
- e) **Wallet Service** (port 8085) — Now with detailed payout analytics via MongoDB

Deliverables: 15 new features (3 per service), JWT authentication on all endpoints, Redis caching on read endpoints, Docker Compose with 6 database containers, continued Git workflow from M1.

Important: This is **not** a microservices milestone. Your services still share a single PostgreSQL database. Inter-service communication comes in Milestone 3. The NoSQL databases are **additive specialized tools**, not replacements for PostgreSQL.

M1 seed data remains valid (additive modifications only): M2 does not alter any M1 PostgreSQL entity schema. The two M1 modifications M2 introduces are all **additive**:

- a) **BCrypt password hashing** on the User entity's existing password field (see Section 4.1) — the column type is unchanged; only how values are stored changes.
- b) **Additive platformFee key** in Payout.transactionDetails JSONB (see Section 4.6). M1's S5-F4 is updated to write this key on new Payouts; existing pre-M2 Payouts without the key default to 0.10 * amount (10% fallback) inside S5-F10's aggregation.

All existing M1 seed rows, tests, and fixtures continue to work after you apply the M1 modifications in Section 4.

Config format (properties → yaml): M1 used application.properties per service. In M2 you must migrate each service's configuration to application.yml. The auto-grader expects YAML format for M2. Reference fragments are provided in Section 6.5.

2 Git Workflow — Feature Development

The same workflow from Milestone 1 continues. Feature IDs are now **F10–F12** per service:

- a) Start from main: `git checkout main && git pull origin main`
- b) Create feature branch: `git checkout -b feat/<service>/S{n}-F{m}/<YOUR-STUDENT-ID>`
- c) Implement: repository methods → service logic → controller endpoint → test via Postman.
- d) Commit: `git commit -m "feat(<service-name>): <description> (<YOUR-STUDENT-ID>)"`
- e) Push and create a Pull Request on GitHub.
- f) At least 1 teammate reviews and approves.
- g) **Merge into main using a regular merge commit (“Create a merge commit” on GitHub). Do NOT use squash merge.** A regular merge preserves the branch name in the merge commit message, which the auto-grader uses to verify who implemented which feature.
- h) **Do NOT delete the feature branch after merging.** The auto-grader will verify that each feature has a correctly-named branch containing the student ID. If you delete the branch, the grader will not be able to verify it and this may result in deductions in your grade.

Example branches for M2 features:

- `feat/user/S1-F10/55-8078` — Register User
- `feat/job/S2-F10/55-8079` — Full-Text Job Search
- `feat/proposal/S3-F11/55-8080` — Record Freelancer-Job Interaction
- `feat/contract/S4-F12/55-8081` — Contract Milestone Timeline
- `feat/wallet/S5-F12/55-8082` — Milestone-Based Payout Reversal

Branch naming convention.

Format: `<type>/<scope>/<descriptor>/<studentID>`

- **type** — one of: feat, fix, hotfix, refactor, docs, test, chore, perf.
- **scope** — the service or area being touched: user, job, proposal, contract, wallet (for per-service work), or m1, cc, infra (for cross-cutting work).
- **descriptor** — the stable ID when the change maps to one (e.g., S3-F11, MOD-3, CC-2), or a short kebab-case otherwise (e.g., payout-rounding-bug).
- **studentID** — your numeric student ID, always the last segment; used by the grader.

Examples:

- `feat/proposal/S3-F11/55-8080` — new feature (M2)
- `feat/m1/MOD-3/55-8080` — M1 amendment

- feat/cc/CC-5/55-8080 — cross-cutting requirement
- fix/wallet/payout-amount-rounding/55-8080 — bug fix
- hotfix/user/token-expiry-leak/55-8080 — urgent fix on an already-merged branch
- refactor/job/extract-search-service/55-8080 — internal cleanup, no behavior change

Commit message convention (Conventional Commits).

Format: <type>(<scope>): <subject> (<studentID>)

- **type** — same set as branch.
- **scope** — same as branch.
- **subject** — imperative mood (“add”, not “added”); no trailing period; keep under 72 characters.
- **studentID** — in parentheses at the end, so every commit is attributable.

Examples:

- feat(proposal): S3-F11 record freelancer-job ordering pattern (55-8080)
- feat(m1): MOD-3 apply JWT filter to all M1 endpoints (55-8080)
- feat(cc): CC-2 role management endpoint (55-8080)
- fix(wallet): correct payout amount rounding (55-8080)
- refactor(user): extract token validation to helper (55-8080)
- test(job): add integration test for S2-F11 auto-index (55-8080)
- docs(cc): clarify CC-5 healthcheck timings (55-8080)
- chore(infra): bump postgres image to 17.4 (55-8080)

Type quick reference.

Type	When to use
feat	New feature, amendment, cross-cutting requirement, or design-pattern implementation.
fix	Bug fix in existing code (wrong behavior, wrong value, wrong status code).
hotfix	Urgent fix on an already-merged branch — use sparingly.
refactor	Code change that neither fixes a bug nor adds a feature (renaming, extracting methods).
docs	Documentation-only change (README, comments, this spec).
test	Tests only (added or fixed).
chore	Build, dependency, or tooling change with no runtime impact.
perf	Performance improvement without behavior change.

Rules:

- One branch → one PR → one work unit. Don’t bundle unrelated changes.
- Always end the branch name and commit subject with your student ID in the exact format shown above — the auto-grader matches on it.
- Prefer the stable IDs (S<n>-F<m>, MOD-<n>, CC-<n>) as the descriptor when the change maps to one. Use a kebab-case slug only for ad-hoc fixes or refactors with no predefined ID.
- Where a design pattern (DP-1..DP-7) is implemented across multiple branches, cite the DP ID in each commit message (e.g., “implements DP-2 Observer”).

Important: Any team member who has no commits in the repository, or whose commits cannot be matched to any feature branch, will receive a **ZERO** as a result.

Important Note: If any member missed any of the git checks (like pushed some change directly on the main and not creating a PR for it) or any other thing, it will get reflected on everyone in the team and not that member only. Same goes for any test case or any check in the entire project.

3 Design Patterns

M2 requires you to apply **7 design patterns** from Lab 7 at specific, well-defined locations in your codebase.

3.1 The 7 Required Patterns

#	Pattern	Category	Where Applied
1	Strategy	Behavioral	S5-F12 payout reversal logic
2	Observer	Behavioral	MongoDB event logging across all services
3	Chain of Responsibility	Behavioral	JWT authentication filter chain
4	Builder	Creational	Dashboard and analytics DTOs
5	Singleton	Creational	JwtService / ConfigurationManager
6	Factory	Creational	MongoDB event object creation
7	Adapter	Structural	NoSQL query result → DTO conversion

Distribution: 3 Creational, 1 Structural, 3 Behavioral — balanced across all 3 pattern categories.

3.2 Strategy Pattern — Payout Reversal Logic (S5-F12)

Purpose: The S5-F12 milestone-based payout reversal feature has multiple business rules that must each be encapsulated as a separate strategy (full payout reversal when the contract is terminated, milestone-only reversal when only specific milestones are disputed, no-reversal when the 30-day reversal window has expired). Replace if-else chains in the service with a Strategy selector.

Structure:

- RefundStrategy interface with calculateRefund(Payout payout, RefundRequest request) method returning a RefundResult (amount + reason code)
- **3 concrete strategies (required):**
 - FullPayoutReversalStrategy — returns the full payout amount when reversalScope=FULL and the payout is within the 30-day reversal window from createdAt
 - MilestoneReversalStrategy — returns the sum of ProposalMilestone.amount values whose status is NOT in {COMPLETED, APPROVED} when reversalScope=MILESTONE_ONLY and the payout is within the 30-day window. Milestones are resolved from the payout's contract via the cross-service native SQL pattern (contract → proposalId → proposal_milestones)
 - NoReversalStrategy — returns zero and a “reversal window expired” reason code when the payout is older than 30 days from createdAt
- A RefundStrategySelector (not the service itself) chooses which strategy to instantiate based on reversalScope and payout age. The service only calls selector.select(payout, request).calculateRefund(payout, request).

Test scenario:

- a) Via reflection, assert that the class RefundStrategy exists and is an interface with exactly one abstract method whose name is calculateRefund.

- b) Via reflection, assert that FullPayoutReversalStrategy, MilestoneReversalStrategy, and NoReversalStrategy all exist and implement RefundStrategy.
- c) Via reflection, assert that a class named RefundStrategySelector (or RefundStrategyFactory) exists with a select-like method returning RefundStrategy.
- d) Call S5-F12 with a payout within the 30-day window and reversalScope=FULL → verify the audit trail in MongoDB contains the strategy name FullPayoutReversalStrategy (or equivalent identifier stored with the event).
- e) Call S5-F12 with a payout within the window and reversalScope=MILESTONE_ONLY → audit trail records MilestoneReversalStrategy.
- f) Call S5-F12 with a payout older than 30 days → returns 400, audit trail records NoReversalStrategy.
- g) Inspect the wallet-service source: the refund service method must not contain if (reversalScope == FULL) or equivalent branching.

3.3 Observer Pattern — Event Logging

Purpose: When entity state changes (user registers, proposal accepted, payout refunded, etc.), events must be logged to MongoDB. Observer decouples logging from business logic.

Structure:

- EntityObserver interface with onEvent(String eventType, Object payload) method
- Concrete observer: MongoEventLogger that writes events to the appropriate MongoDB collection
- Subjects (services) maintain an observer list and call notifyObservers(eventType, payload) on state changes. Provide register(observer) / unregister(observer) methods on each subject.
- Each of the 5 services owns its own per-service MongoEventLogger instance; observer registration is not shared across services.

Failure policy Catching Exception for Mongo related features (Soft Dependency): The MongoEventLogger must catch any Mongo exception, log it at WARN level (log.warn()), and **not rethrow**. The upstream Postgres transaction must **not** be rolled back on a Mongo write failure.

Spring vs classical GoF: Spring provides a built-in Observer mechanism via ApplicationEventPublisher + @EventListener. For M2, **you must implement the classical GoF Observer** (explicit interface + register/notify methods) for learning purposes. No method in any of the 5 services annotated with @EventListener may write to MongoDB — all MongoDB event writes must flow through your classical Observer chain. You may, of course, still use Spring for everything else in your services.

Where applied:

- S1 register/login/role-change → observers log AuthEvent to auth_events
- S2 job updates → observers log to job_events
- S3 proposal lifecycle changes → observers log to proposal_events
- S4 contract milestone tracking → observers log to contract_events
- S5 payout reversal processed → observers log to payout_audit_trail
- **Also applied to your M1 write endpoints** (see Section 4)

Test scenario:

- a) Via reflection, assert that EntityObserver exists as an interface with an onEvent(String, Object) method.
- b) Via reflection, assert MongoEventListener class exists and implements EntityObserver.
- c) Via static analysis (class-file or source scan), assert that no method in the User/Job/Proposal/Contract/Wallet services annotated with @EventListener writes to MongoDB at all — all event-logging to MongoDB must flow through the classical GoF Observer chain.
- d) Register a fresh user via POST /api/auth/register → verify exactly one REGISTERED document exists in auth_events collection with matching userId and recent timestamp.
- e) Log in as that user → verify a LOGGED_IN document is added.
- f) Trigger an M1 write (e.g., PUT /api/users/{id}/preferences from S1-F2) → verify the corresponding event document is written to MongoDB (confirms the retrofit path).
- g) Unregister all observers from a service in a unit test and repeat the write → no MongoDB event should appear (proves the logging path goes through the observer chain, not a direct Mongo call).

3.4 Chain of Responsibility — JWT Filter Chain

Purpose: JWT authentication has multiple sequential steps (extract token, verify signature, load user, check role). Each step is a separate handler in a chain.

Structure:

- AuthHandler abstract class (or interface) with handle(AuthContext ctx) and setNext(AuthHandler next) methods, where AuthContext carries the HTTP request, extracted token, authenticated user, and required role
- Concrete handlers (in order): TokenExtractionHandler (401 if missing), SignatureValidationHandler (401 if invalid/expired), UserLoaderHandler (401 if user not found in PG), RoleAuthorizationHandler (403 if insufficient role for the endpoint)
- Each handler does its job and passes to the next; returns an error result immediately if any check fails

Spring Security integration: Spring Security's SecurityFilterChain is itself a chain of responsibility at the framework level. You **must not** replace Spring Security's filter chain. Instead, build your custom AuthHandler chain **inside** your JwtAuthenticationFilter.doFilterInternal() method:

- Spring Security's filter chain invokes your JwtAuthenticationFilter once per request
- Inside doFilterInternal, construct (or inject) the head of your custom AuthHandler chain and call head.handle(new AuthContext(request))
- If any handler in your chain fails, write the appropriate status code (401 or 403) to the response and short-circuit by **not** calling filterChain.doFilter(...)
- If all handlers succeed, populate the Spring Security context (SecurityContextHolder.getContext().setAuthentication(...)) and call filterChain.doFilter(...)

Test scenario:

- a) Via reflection, assert AuthHandler class/interface exists with setNext(AuthHandler) and handle(...) methods.
- b) Via reflection, assert at least 3 concrete subclasses of AuthHandler exist (e.g., TokenExtractionHandler, SignatureValidationHandler, UserLoaderHandler, and optionally RoleAuthorizationHandler).

- c) Call a protected M1 endpoint (e.g., GET /api/users/1) with Authorization header absent → returns 401 (TokenExtractionHandler failed).
- d) Call with Authorization: Bearer invalid.token.here (invalid signature) → returns 401 (SignatureValidationHandler failed).
- e) Delete the user from PG, then call with a valid token issued for that user → returns 401 (UserLoaderHandler failed).
- f) Call PUT /api/users/1/role with a CLIENT token → returns 403 (RoleAuthorizationHandler failed).
- g) Call the same endpoint with an ADMIN token → succeeds (chain passes through).
- h) Inspect JwtAuthenticationFilter.doFilterInternal() source: the grader parses it and confirms that the method body invokes the head of the AuthHandler chain rather than duplicating the extraction/validation/authorization logic inline.

3.5 Builder Pattern — Dashboard DTOs

Purpose: Dashboard and analytics DTOs have 5–10+ fields. Builder improves readability and allows optional fields.

Structure:

- Static inner Builder class on each dashboard DTO (or external Builder class)
- builder() static method to start construction
- Fluent setters returning this
- build() method returning the DTO instance

Builder with Java records: The project convention is “Java records for DTOs.” Records do not naturally support Builder (the canonical constructor takes all fields positionally). Two options, both accepted by the grader:

- **Convert to class:** Turn the record into a class with a static inner Builder.
- **External builder on the record:** Keep the record and write an external <DtoName>Builder class whose build() method constructs the record via its canonical constructor.

Where applied in M2: S2-F12 JobDashboardDTO, S3-F10 ProposalAnalyticsDashboardDTO, S4-F10 ContractAnalyticsDTO, S5-F10 CategoryRevenueDTO.

Freelance M1 retrofit scope: Apply Builder to M1 DTO-returning features whose DTOs have 5+ fields. These DTOs for Freelance specifically are:

- **S1-F3, S1-F6, S1-F8, S1-F9:** Builder required (all return DTOs with 5+ fields)
- **S2-F3, S2-F6, S2-F9:** Builder required
- **S2-F8 (Verify Job Attachment): No Builder retrofit** — this feature returns an entity, not a DTO
- **S3-F3 (Platform Fee Estimate — POST estimate), S3-F6, S3-F9:** Builder required
- **S3-F8 (Add Milestones to Proposal): No Builder retrofit** — returns entities
- **S4-F3, S4-F6, S4-F8, S4-F9:** Builder required
- **S5-F3, S5-F6, S5-F8, S5-F9:** Builder required

Test scenario:

- a) Via reflection, iterate every DTO class returned by an M2 dashboard feature: assert each has an accessible static builder() method, that chaining setters returns the Builder, and that build() returns the correct DTO type.
- b) Via reflection, iterate the in-scope Freelance M1 DTOs (UserContractSummaryDTO, TopFreelancerDTO, JobProposalSummaryDTO, etc.): each must have a Builder.
- c) In an integration test, call S2-F12 GET /api/jobs/{id}/dashboard → verify the response is correctly populated (proves Builder is used to construct the DTO before serialization).
- d) Call M1 S1-F3 GET /api/users/{id}/contract-summary → verify it still returns the expected fields (proves the retrofit did not break M1 behavior).
- e) Compile-time or static analysis check: confirm S2-F8 and S3-F8 do **not** use Builder (they return entities, not DTOs).

3.6 Singleton Pattern — JwtConfigurationManager

Purpose: A class holding shared, immutable JWT configuration (secret key, expiration, algorithm) should have exactly one instance.

Structure:

- Private constructor
- Static getInstance() method with thread-safe initialization (double-checked locking or eager init)
- Single private static field holds the instance

Loading config in a non-Spring class: Since JwtConfigurationManager is not a Spring bean, it cannot use @Value or @ConfigurationProperties. Use one of the following accepted approaches:

- **(Recommended) Environment variables with fallback defaults:** Inside the private constructor, read System.getenv("JWT_SECRET"), System.getenv("JWT_EXPIRATION_MS"), etc., with sensible fallback defaults for local dev. Docker Compose already sets the required env vars.
- **Singleton-bridge pattern:** A Spring @Configuration bean reads application.yml via @Value, then pushes the values into the singleton via a static setter (e.g., JwtConfigurationManager.initConfig(secret, expirationMs)) during @PostConstruct. The singleton then serves those values through its instance methods.

Both approaches are accepted by the grader as long as the class itself is not a Spring bean and getInstance() returns a consistent singleton reference.

Spring vs classical GoF (important): Spring's @Component/@Service beans are singletons by default, but they are managed by the Spring container, not by the GoF pattern. For M2 you must have **exactly one** class implemented as a classical GoF Singleton:

- **Implement** JwtConfigurationManager **as a classical Singleton** (private constructor + getInstance()). This class is not a Spring bean. Do not annotate it with @Component.
- JwtService **remains a Spring** @Service. It obtains JWT config via JwtConfigurationManager.getInstance() rather than via @Autowired.
- All other infrastructure classes (controllers, services, repositories) remain Spring-managed.

Test scenario:

- a) Via reflection, assert JwtConfigurationManager exists and has exactly one constructor declared with private access.

- b) Via reflection, assert `getInstance()` is a public static method returning `JwtConfigurationManager`.
- c) Call `getInstance()` twice in the test and assert `ref1 == ref2` (reference equality, not just `.equals()`).
- d) Spawn 10 parallel threads each calling `getInstance()` → all must return the same reference (thread-safety under contention).
- e) Via reflection, assert the class is NOT annotated with `@Component`, `@Service`, `@Configuration`, or any other Spring stereotype.
- f) In an integration test, verify that `JwtService` (a Spring bean) correctly reads the configured secret via `JwtConfigurationManager.getInstance()` — issue and validate a token and confirm the flow works.

3.7 Factory Pattern — Event Creation

Purpose: MongoDB events have different types across services (`AuthEvent`, `JobEvent`, `ProposalEvent`, `ContractEvent`, `PayoutAuditEvent`). A `EventFactory` creates the right concrete event type based on input.

Structure:

- A common `MongoEvent` interface that all 5 concrete event classes implement, with methods `getId()`, `getTimestamp()`, `getAction()`, `getDetails()`. (See Section 7.1 for the common interface definition.)
- `EventFactory` class with `createEvent(EventType type, Map<String, Object> params)` method
- The factory dispatches on type (an enum with values `AUTH`, `JOB`, `PROPOSAL`, `CONTRACT`, `PAYOUT_AUDIT`) and returns the appropriate concrete class, typed as `MongoEvent`

How Observer and Factory compose: On an M1 write completing, the service calls `notifyObservers(eventType, payload)`. The `MongoEventLogger` observer receives the call, asks the `EventFactory` for the right event subtype, and persists the event to MongoDB via Spring Data.

Test scenario:

- a) Via reflection, assert `MongoEvent` interface exists with methods `getId`, `getTimestamp`, `getAction`, `getDetails`.
- b) Via reflection, assert all 5 event classes (`AuthEvent`, `JobEvent`, `ProposalEvent`, `ContractEvent`, `PayoutAuditEvent`) implement `MongoEvent`.
- c) Via reflection, assert `EventFactory` exists with a `createEvent(EventType, Map<String, Object>)` method.
- d) In a unit test, call `EventFactory.createEvent(AUTH, params)` → assert the returned object is assignable to `AuthEvent` and its fields match the params.
- e) Repeat for each of the 5 `EventType` values.
- f) Call `EventFactory.createEvent(PAYOUT_AUDIT, ...)` → the returned `PayoutAuditEvent` must expose method and amount fields (service-specific fields added on top of the common interface).
- g) Integration test: register a user via `POST /api/auth/register` → inspect the resulting MongoDB document in `auth_events`. Its type and fields must match what `EventFactory.createEvent(AUTH, ...)` would produce (proves services go through the factory).
- h) Source-scan check: verify no service class contains `new AuthEvent(...)`, `new JobEvent(...)`, etc. — all event construction must go through the factory.

3.8 Adapter Pattern — NoSQL Result → DTO Conversion

Purpose: Results from different NoSQL databases have different raw shapes (MongoDB Document, Elasticsearch SearchHit, Neo4j Record, Cassandra Row). Adapters convert each raw source to the specific DTO the service returns.

Structure:

- Each adapter has a single `adapt(source) → targetDto` method. There is **no** universal “Entity-Dto” base type — each adapter converts to its specific domain DTO.
- One adapter per NoSQL source used in the service:
 - S1: MongoDocumentAdapter (for AuthEvent documents)
 - S2: MongoDocumentAdapter, ElasticsearchHitAdapter
 - S3: MongoDocumentAdapter, Neo4jRecordAdapter
 - S4: MongoDocumentAdapter, CassandraRowAdapter
 - S5: MongoDocumentAdapter

M1 retrofit scope (conditional): In M1, some features that return DTOs from native SQL used `Object[]` row projections while others used JPQL constructor expressions or `@Query` DTO projections. The Adapter retrofit is **conditional**:

- **If** your M1 feature uses `Object[]` (S1-F3 in the M1 spec explicitly mandates this) → wrap the mapping in an `ObjectArrayDtoAdapter` class rather than doing inline mapping in the service.
- **If** your M1 feature uses JPQL constructor expressions or DTO projections → no Adapter needed; the existing code satisfies the pattern by construction.

Test scenario:

- a) Via reflection, assert that for each NoSQL source used by each service, a corresponding adapter class exists: `MongoDocumentAdapter` in all services; `ElasticsearchHitAdapter` in job-service; `Neo4jRecordAdapter` in proposal-service; `CassandraRowAdapter` in contract-service.
- b) Via reflection, assert each adapter has an `adapt(...)` method whose return type matches the service’s domain DTO.
- c) In a unit test, pass a mock MongoDB Document to `MongoDocumentAdapter.adapt(...)` → verify the returned DTO has fields populated from the document’s keys.
- d) In a unit test for job-service: pass an `Elasticsearch SearchHit` to `ElasticsearchHitAdapter.adapt(...)` → verify the returned job DTO.
- e) For S1-F3 (the only M1 feature that explicitly mandates `Object[]`): via reflection or source scan, verify an `ObjectArrayDtoAdapter` (or similar) class exists and is used to convert the native SQL `Object[]` rows into `UserContractSummaryDTO`. Invoke S1-F3 in integration and verify the output is correct.
- f) For any other in-scope M1 F3/F6/F9 feature that the team chose to implement with `Object[]`: check the corresponding adapter exists. Features implemented via JPQL constructor expressions or DTO projections are **exempt** from this check.

3.9 Grading Summary

Each pattern contributes roughly equally to the design patterns portion of the M2 grade. The grader performs:

- **Static analysis** — reflective inspection of class structure (interfaces, methods, constructors)
- **Behavioral verification** — integration tests that exercise the patterns indirectly through M2 features

4 M1 Modifications Required

Before you can build M2 features, you must **modify your existing M1 codebase** in specific ways. These modifications are prerequisites — M2 features and cross-cutting requirements depend on them.

4.1 User Entity — Password Hashing

Your M1 User entity already has a password field. In M1 it was stored as plaintext (or loosely). In M2:

- All new user registrations must hash passwords with **BCrypt** before saving
- Existing M1 test users/seeds should be re-created with hashed passwords (or hash them on first login)
- The password field must **never** be returned in any API response (including existing M1 GET /api/users/{id} endpoint — filter it out in the DTO)

Test scenario:

- a) Register a user via POST /api/auth/register with password "securePassword123".
- b) Query the users table directly and assert the stored password column starts with "\$2a\$", "\$2b\$", or "\$2y\$" (BCrypt hash prefix) and its length is 60 characters.
- c) Assert the stored password is **not** equal to "securePassword123".
- d) POST /api/auth/login with the same credentials → succeeds (BCrypt verify works).
- e) Call GET /api/users/{id} (M1 endpoint) with a valid token → inspect the JSON response body: the password field must be absent or explicitly null. Same for any other endpoint that returns a User (e.g., S1-F1 search, S1-F3 summary).
- f) Re-seed the database using the seed mechanism → assert every seeded user's password column is a BCrypt hash, not plaintext.

4.2 User Entity — Role Values (Additive)

Your M1 User entity already has a role ENUM with values FREELANCER, CLIENT, and ADMIN. **Do not remove or rename existing M1 role values.** M2 modifies this only additively:

- **Keep** FREELANCER, CLIENT, **and** ADMIN as they are
- **Default role on registration:** CLIENT (Freelance's M1 default). The register payload may optionally include a role field whose value is either CLIENT or FREELANCER — if provided, honour it. If the payload contains role=ADMIN, the server must silently ignore that value and still create a CLIENT. ADMIN is never assigned at registration.
- **ADMIN assignment:** Only via the role management endpoint (PUT /api/users/{id}/role) which itself requires ADMIN. Seed at least one ADMIN user via your M1 seed mechanism so role management can be tested.

JWT role claim: The token carries the user's role. Authorization logic:

- **USER-level endpoints:** Accept any authenticated user (CLIENT, FREELANCER, or ADMIN) — this is the default for M2 features.
- **ADMIN-only endpoints:** Require role == ADMIN. Only the role management endpoint uses this.

Test scenario:

- a) Query PostgreSQL metadata (information_schema.columns or the Postgres enum pg_enum) to confirm the role ENUM type contains FREELANCER, CLIENT, and ADMIN values.
- b) Register a fresh user via POST /api/auth/register with valid data and no explicit role field → query the users table and assert the new row's role column is CLIENT.
- c) Decode the returned JWT token → assert the role claim is CLIENT.
- d) Register a fresh user with "role": "FREELANCER" in the body → the new row's role column is FREELANCER.
- e) Log in as a seeded ADMIN user → decode token and assert role claim is ADMIN.
- f) Attempt to register with a request body containing "role": "ADMIN" (a bad-faith attempt) → the server must ignore that field and still create a CLIENT.
- g) With a CLIENT token, call PUT /api/users/{id}/role → 403 Forbidden.
- h) With the ADMIN token, call PUT /api/users/{id}/role with body {"role":"FREELANCER"} → 200, the target user's role changes.
- i) With a CLIENT token, call any regular M2 feature (e.g., S2-F10 search) → succeeds (USER-level endpoints accept every role).

4.3 Existing M1 Endpoints — JWT Authentication

All existing M1 endpoints (45 feature endpoints + CRUD endpoints on all entities) must now require a valid JWT token, except:

- POST /api/auth/register (new in M2)
- POST /api/auth/login (new in M2)
- Health check endpoints

Test scenario:

- a) Call POST /api/auth/register (public) → succeeds without any Authorization header.
- b) Call POST /api/auth/login (public) → succeeds without a token.
- c) Call any M1 feature endpoint (e.g., GET /api/jobs/search, M1 S2-F1) without token → returns 401.
- d) Call the same endpoint with Authorization: Bearer invalid → returns 401.
- e) Call the same endpoint with a valid token → succeeds with 200 and the expected M1 response shape.
- f) Call every CRUD endpoint (GET /api/users, GET /api/users/{id}, POST /api/users, PUT /api/users/{id}, DELETE /api/users/{id}, and the equivalents on all 10 Freelance entities) without token → each returns 401. With a valid token → each returns the expected M1 status code.
- g) Health check endpoints (e.g., GET /actuator/health) return 200 without a token.
- h) Confirm by grepping the test results that all 45 M1 feature tests + CRUD tests pass under M2 auth.

4.4 Existing M1 Read Endpoints — Redis Caching (Explicit Enumeration)

All M1 read-heavy endpoints must be cached in Redis. Below is the complete enumeration — nothing is left open.

4.4.1 M1 Feature GET Endpoints That Must Be Cached (Freelance: 27 total)

The canonical M1 pattern is “F2, F4, F7 are writes, the other 6 are reads,” but **Freelance deviates at three specific features** that are writes despite their F-number:

Service	Feature	Verb	Why it is a write
S2 (Job)	S2-F8 Verify Job Attachment	PUT	Modifies JobAttachment.verified
S3 (Proposal)	S3-F8 Add Milestones to Proposal	POST	Creates new ProposalMilestone rows
S5 (Wallet)	S5-F5 Apply PromoCode to Payout	POST	Creates PayoutPromo join row and mutates Payout

These are **excluded** from the cached-reads list and **included** in the invalidating-writes list (Â§4.4.4). The explicit cached-read enumeration for Freelance is:

Service	Cached GET feature IDs
S1 (User)	F1, F3, F5, F6, F8, F9
S2 (Job)	F1, F3, F5, F6, F9 (F8 is a write)
S3 (Proposal)	F1, F3, F5, F6, F9 (F8 is a write; F3 is a POST estimate—see not
S4 (Contract)	F1, F3, F5, F6, F8, F9
S5 (Wallet)	F1, F3, F6, F8, F9 (F5 is a write)
Total cached feature endpoints (Freelance)	

Note on S3-F3 (Get Platform Fee Estimate): In the M1 spec S3-F3 is a POST /api/proposals/estimate that computes a platform-fee estimate without persisting. It is semantically read-only despite being POST. Cache it by request-body hash (TTL 5 min). If you did not implement S3-F3 as a POST, follow the M1 spec verb for your service.

TTLs by feature type: F1 = 5 min (search), F3 = 10 min (DTO), F5 = 5 min (JSONB query), F6 = 10 min (report), F8 = 15 min (relationship DTO), F9 = 10 min (combined).

4.4.2 CRUD Baseline Endpoints That Must Be Cached

For every entity, **only** the get-by-ID endpoint is cached. List endpoints are **not** cached — they hit PostgreSQL on every request.

Endpoint	Cached?	TTL
GET /api/<entity> (list all)	No	—
GET /api/<entity>/<id> (get by ID)	Yes	15 min

Freelance entities (10): user, user-skill, job, job-attachment, proposal, proposal-milestone, contract, payout, promo-code, payout-promo. **10 GET-by-ID endpoints must be cached.**

4.4.3 Total

27 M1 feature endpoints + 10 M1 CRUD GET-by-ID endpoints = 37 M1 endpoints must be cached for Freelance. Plus the M2 feature GET endpoints.

4.4.4 Endpoints That Must Invalidate Caches (Writes)

Writes invalidate the detail cache of the specifically affected entity ID, plus feature-result caches whose output includes that entity. Since list endpoints are not cached, there is no list cache to invalidate.

Freelance M1 feature writes (18 total):

- S1 writes: F2, F4, F7 (3)
- S2 writes: F2, F4, F7, **F8 (Verify Job Attachment)** (4)
- S3 writes: F2, F4, F7, **F8 (Add Milestones to Proposal)** (4)
- S4 writes: F2, F4, F7 (3)
- S5 writes: F2, F4, **F5 (Apply PromoCode to Payout)**, F7 (4)

M1 CRUD writes (30): POST, PUT, DELETE on each of the 10 Freelance entities.

- POST /api/<entity> (create) — nothing to invalidate (new entity has no cached detail yet)
- PUT /api/<entity>/<id> (update) — invalidate detail cache for that specific ID + any feature caches referencing that entity
- DELETE /api/<entity>/<id> — invalidate detail cache for that specific ID + any feature caches referencing that entity

Total Freelance M1 endpoints that invalidate: 18 feature writes + 30 CRUD writes = 48.

M2 writes that also invalidate M2 feature caches: In addition to the M1 writes above, the following M2 operations must invalidate feature-result caches so that read-side views see the change before the TTL expires:

- CC-2 Role management (PUT /api/users/<id>/role) — must invalidate user-service::user::<id> (detail) and all keys matching user-service::S1-F12::* for that userId (so the ROLE_CHANGED event appears in the activity feed immediately).
- Any M1 write that creates/updates a Proposal, ProposalMilestone, or Contract referencing a specific jobId — must invalidate job-service::S2-F12::{jobId} so the Job Market Dashboard reflects the new activity.
- Any M1 write that creates/updates a Proposal — must invalidate proposal-service::S3-F10::* (the Proposal Analytics dashboard aggregates over all proposals).
- Any M1 write that creates/updates a Contract — must invalidate contract-service::S4-F10::*.
- Any M1 write that creates/updates a Payout (including S5-F4 Process Payout, S5-F2 Refund, M2 S5-F12 Milestone Reversal) — must invalidate wallet-service::S5-F10::* and wallet-service::S5-F11::*.

NoSQL-writer → cached-reader invalidation (required): Several M2 features write to a NoSQL store as their primary action (no PG row mutates), but downstream M2 features cache reads of that same NoSQL store. The PG-write rules above do **not** cover these paths, so add these explicitly:

- **S4-F11** (POST /api/contracts/<id>/milestones/track) writes a new milestone row to Cassandra and emits a MILESTONE_TRACKED event to MongoDB contract_events — must invalidate contract-service::S4-F12::{contractId} (the cached milestone timeline for that contract, 5 min TTL) so S4-F12 reflects the new milestone event immediately. Must also invalidate contract-service::S4-F10::* (the contract analytics dashboard may aggregate over milestone-event counts).

- **S3-F11** (POST /api/proposals/{proposalId}/record-interaction) writes a new PROPOSED_TO relationship (or increments an existing one) in Neo4j — must invalidate proposal-service::S3-F12:* (the cached recommendation results for all freelancers, 5 min TTL) because the recommendation graph changed. Use wildcard deletion rather than narrow targeting: a new edge between freelancer F and job J can alter recommendations for every freelancer who shares a job with F, not just F.
- **S2-F11** (POST /api/jobs/{id}/index) **and** every Job CRUD write that triggers the auto-indexing retrofit (see §10 S2-F11 step e) — must invalidate job-service::S2-F10:* (the cached full-text search results, 5 min TTL). Job title / description / status changes alter ES relevance ranking, so stale cache can return deleted or renamed jobs.
- **S5-F10 / S5-F11 analytics writes** — whenever any Observer writes a **data-mutating** event to payout_audit_trail (specifically CREATED, COMPLETED, FAILED, REFUNDED, REFUND_DENIED, PROMO_APPLIED), the wallet-service analytics caches wallet-service::S5-F10:* and wallet-service::S5-F11:* must be invalidated. This covers both the normal-processing path (S5-F4, S5-F2, S5-F5, S5-F12 success branch) and the strategy-denial path (S5-F12 NoReversalStrategy branch logs REFUND_DENIED without any PG mutation — the PG-write rules above miss this case, so the analytics caches would otherwise serve stale data until TTL expiry). Similarly, any Observer write of a data-mutating action to job_events, proposal_events, or contract_events invalidates the corresponding S2-F12 / S3-F10 / S4-F10 analytics keys.
- **Pure observability actions do NOT invalidate caches** — specifically ANALYTICS_VIEWED and DASHBOARD_VIEWED. These actions are written by the four dashboard endpoints (S2-F12, S3-F10, S4-F10, S5-F10) on *every* invocation (including cache hits; see §10) purely for audit-trail / usage-telemetry purposes. They do not change what the analytics queries return. Triggering wildcard invalidation on them would create a self-defeating cycle: every cached dashboard call would log an observability event, which would wildcard-invalidate its own key, guaranteeing that the next call is always a miss and the cache never serves a hit. Exclude both ANALYTICS_VIEWED and DASHBOARD_VIEWED from the invalidation trigger in the Observer (match on the action field before invalidating).

These M2 invalidation rules compose with the M1 rules above; implement both.

4.4.5 Cache Key Convention

- Entity detail: <service>::<entity>::<id> (e.g., user-service::user::42)
- Feature result: <service>::<featureId>::<param-hash> (e.g., user-service::S1-F3::42)

4.4.6 Cache Invalidation Strategy (Wildcard Deletion)

Computing the exact set of feature-result caches affected by an entity change is infeasible without cache tags or reverse indexes. Use **wildcard key deletion** instead:

- When entity X with id N changes, delete the key <service>::X::N (entity detail).
- Also delete all feature-result keys matching the pattern <service>::S{n}-F{m}::* for every feature S{n}-F{m} whose output *may* include entity X (conservatively identified at design time).
- This over-invalidates (clears some entries that did not actually change). That is acceptable for M2 — correctness beats cache-hit optimality.

Use the Redis SCAN + DEL combination or KEYS-UNLINK (for small caches) to implement wildcard deletion.

Test scenario (entire §4.4 caching retrofit):

- a) For each of the 27 cached M1 feature GET endpoints listed in §4.4.1: call it twice with identical parameters and a valid token. Inspect Redis between calls (e.g., via redis-cli KEYS

'user-service::') and assert a matching key exists. Measure the second call's latency — it must be less than the first (cache hit).

- For each of the 10 CRUD GET /api/<entity>/<id> endpoints: call twice and verify a key <service>::<entity>::<id> exists in Redis after the first call.
- Call GET /api/jobs (list) twice → assert **no** cache key was created for the list (list endpoints are NOT cached).
- Fetch job ID 5 (cache it), then PUT /api/jobs/5 with an update → verify the key job-service::job::5 is removed from Redis. Next GET recomputes from PG.
- Fetch the S1-F3 DTO for user 10 (cache it), then update user 10 via PUT /api/users/10/preferences (S1-F2) → verify every key matching user-service::S1-F3::* is removed (wildcard invalidation).
- Delete job 7 via DELETE /api/jobs/7 → the cached job-service::job::7 must disappear.
- Disable Redis (e.g., stop the container) and retry a cached endpoint → it still returns correct data from PG (soft-dependency graceful degradation).
- Verify TTLS: fetch an S1-F1 search result (5 min TTL), wait 5m+, fetch again → cache was auto-evicted and the call recomputes.

4.5 Design Pattern Retrofits to M1

The 7 patterns from Section 3 apply to M2 code **and** to M1 retrofits as follows:

Pattern	Applies to M1?	Which M1 Freelance features	Notes
Strategy	No		New M2 code only (S5-F12 milestone-based reversal) — do NOT force onto any M1 method
Observer	Yes (universal)	All M1 write endpoints: F2, F4, F7 per service + S2-F8, S3-F8, S5-F3 (Freelance-specific writes) + all CRUD writes	State change → notifyObservers → MongoDB event log
Chain of Responsibility	Yes (universal)	All M1 endpoints via the JWT filter chain	Inside JwtAuthenticationFilter (see Section 3.4)
Builder	Yes	S1-F3, S1-F6, S1-F8, S1-F9, S2-F3, S2-F6, S2-F9, S3-F3, S3-F6, S3-F9, S4-F3, S4-F6, S4-F8, S4-F9, S5-F3, S5-F6, S5-F8, S5-F9 (DTO-returning features with 3+ fields)	Excludes S2-F8 and S3-F8 for Freelance — they return entities, not DTOs
Singleton	No		New M2 code only (JwtConfigurationManager)
Factory	Yes (indirect)	M1 write endpoints trigger Observer; Observer calls EventFactory	EventFactory creates the right MongoEvent subtype
Adapter	Conditional	Any M1 feature that uses Object[] native SQL projection (S1-F3 mandates Object[]); other F3/F6/F9 only if you chose Object[] in M1	If you used JPQL constructor expressions or DTO projections, no Adapter retrofit is required

Important: Strategy Pattern is used **only** in M2 code (S5-F12 milestone-based reversal). Do not force it onto any M1 method.

Composition workflow (Observer + Factory + Adapter):

- A write endpoint (e.g., M1 S1-F2 Update User Preferences) completes its PG update.
- The service calls notifyObservers("USER_UPDATED", userPayload) before returning. The first argument is the **action string** (what happened); it is not the EventType passed to the factory.
- MongoEventLogger observer receives the call. Each of the 5 services instantiates its own MongoEventLogger bound to a **fixed** EventType at construction time (user-service binds AUTH, job-service binds JOB, proposal-service binds PROPOSAL, contract-service binds CONTRACT, wallet-service binds PAYOUT_AUDIT). The logger does **not** infer EventType from the action string.
- The logger builds the factory input: copy the action string into params.put("action", eventType) plus the rest of the payload, then calls EventFactory.createEvent(boundEventType, params).
- The factory returns the matching concrete MongoEvent subtype (e.g., AuthEvent) typed through the common interface.
- The logger persists the event to MongoDB via Spring Data's repository.

Adapter sits separately on **read paths** (converting NoSQL query results or M1 Object[] results to DTOs) — it does not participate in the write-side composition above.

Wallet Service payout_audit_trail population (important): The payout_audit_trail collection powers M2 S5-F11 (GET /api/payouts/analytics/methods). Without writes to this collection, S5-F11 returns empty results. Therefore:

- **M1 S5-F4** (Process Payout for Contract) must write two events to payout_audit_trail: a CREATED event when the Payout row is first inserted, and a COMPLETED event when the payout status transitions to COMPLETED.
- **M1 S5-F2** (Process Refund — the original M1 refund) must write a REFUNDED event.
- **M2 S5-F12** (Process Milestone-Based Payout Reversal) must also write a REFUNDED event with richer details (strategy name, reversal scope).

These are observer-driven writes — use the Observer retrofit (row 2 of the table above) rather than inline Mongo calls in the service methods.

Additional M1 behavior retrofits:

- **S5-F4 gateway-failure simulation (query param):** M1's POST /api/payouts/contract/{contractId} must accept an optional ?simulateFailure=true query parameter. When set, S5-F4 short-circuits to Payout.status = FAILED in PostgreSQL and writes a FAILED event to payout_audit_trail (via Observer) instead of the normal CREATED→COMPLETED flow. When absent or false, S5-F4 behaves identically to its M1 spec. This affordance exists solely so the S5-F11 test scenario can produce failed-payout data; it is not a production flow. See the FAILED row of §7.1.6 PayoutAuditEvent for the full write-conditions table.
- **Job CRUD auto-index to Elasticsearch:** M1's Job CRUD controller must auto-sync to the M2 Elasticsearch index on every write. POST /api/jobs and PUT /api/jobs/{id} must **re-index** the job's search document (as if S2-F11 had been called explicitly); DELETE /api/jobs/{id} must **remove** the document from the index. See §10 S2-F11 step e for the exact expected behavior. Implement via a @PostPersist/@PostUpdate/@PostRemove JPA entity listener or a service-level hook; do **not** inline the ES call into every CRUD controller method.

Test scenario (design pattern retrofits):

- a) Call M1 S1-F2 PUT /api/users/{id}/preferences → verify an event document appears in auth_events (Observer retrofit fired on the write).
- b) Call M1 S2-F8 (Verify Job Attachment) → event in job_events.
- c) Call M1 S3-F8 (Add Milestones to Proposal) → event in proposal_events.
- d) Call M1 S5-F5 (Apply PromoCode to Payout) → event in payout_audit_trail.
- e) Call M1 S5-F4 (Process Payout) → CREATED and COMPLETED events appear in payout_audit_trail.
- f) Call M1 S5-F2 (Process Refund) → REFUNDED event appears.
- g) Call a CRUD write (e.g., POST /api/jobs) → event in job_events.
- h) Call M1 S1-F3 GET /api/users/{id}/contract-summary → response is well-formed; via source scan, confirm the DTO is constructed via its Builder (not via new UserContractSummaryDTO(...) directly inside the service).
- i) Call M1 S1-F3 and confirm that the Object[] result from the native SQL query is converted to UserContractSummaryDTO via an Adapter class (not inline mapping in the service).
- j) Verify that Strategy is NOT used anywhere in M1 service classes (grep for RefundStrategy, Strategy outside the wallet service's S5-F12 package).

4.6 Payout.transactionDetails — Additive platformFee Key

M2 S5-F10 (Platform Fee Analytics by Job Category) requires the Payout’s platform fee to be separable from the net payout amount. M1’s Payout.transactionDetails JSONB schema is {gatewayResponse, accountLastFour, receiptUrl, failureReason} — it does **not** include a platformFee key.

M2 additively extends this JSONB with a new key:

- platformFee — numeric amount (currency units) representing the platform-fee portion of the total payout.

Required M1 code change: Update M1’s S5-F4 (Process Payout) to compute the platform fee when creating the Payout and write it into transactionDetails.platformFee. Derive the platform fee as a flat 10% of the contract’s agreedAmount (Freelance’s default fee model), or from a job-category-level configuration if the team implements one.

Fallback for pre-M2 rows: Existing M1 Payout rows created before this modification will not have a platformFee key. Any feature reading the key (e.g., S5-F10) must treat a missing or null platformFee as $0.10 * \text{payouts.amount}$ (10% of total). This keeps pre-M2 data usable without a DB backfill migration.

4.7 Summary Checklist

Before starting M2 features, ensure your M1 codebase:

- ☐ Password hashing (BCrypt) applied to registration and seed data
- ☐ ADMIN role present; CLIENT preserved as default; FREELANCER preserved
- ☐ At least one ADMIN user seeded
- ☐ All endpoints (except register/login/health) require JWT
- ☐ 27 M1 feature GET endpoints + 10 CRUD GET-by-ID endpoints cached (37 total — see Section 4.4)
- ☐ M1 write endpoints invalidate the correct detail/feature caches
- ☐ M1 write endpoints notify observers (event logging)
- ☐ M1 complex DTOs (F3/F6/F8/F9) use Builder pattern
- ☐ M1 Object[] native SQL results (F3/F6/F9) use Adapter pattern
- ☐ M1 S5-F4 writes platformFee to Payout.transactionDetails JSONB — see Section 4.6
- ☐ All existing M1 tests still pass

5 Authentication & Authorization

5.1 Overview

The User Service (S1) becomes the authentication provider. It exposes register and login endpoints that return JWT tokens. All other services validate tokens independently using a shared secret key — no inter-service calls are needed for authentication.

5.2 JWT Token

- **Algorithm:** HMAC-SHA256
- **Secret:** A shared key configured in each service's application.yml (e.g., jwt.secret). The shared secret must be at least **32 bytes (256 bits)** of entropy when decoded — JJWT's HS256 signer enforces this minimum and throws `WeakKeyException` at construction time on shorter keys. Use `Keys.secretKeyFor(SignatureAlgorithm.HS256)` to generate a compliant key, or configure any 44-character Base64-encoded string of random bytes (44 Base64 chars decode to 32 bytes). Short readable strings like "mySecret123" will fail at the first token-issue attempt.
- **Payload:** The token contains the user's email (as the sub subject claim), the user's numeric `User.id` (as a custom uid claim), the user's role (as a role claim), issued-at (iat), and expiration (exp). Ownership checks in §10 (S1-F12, S3-F12) compare the uid claim directly against the path/query `userId` — no additional PG lookup is required on the hot path.
- **Header format:** Authorization: Bearer <token>
- **Expiration:** Configurable (e.g., 24 hours)

5.3 Roles (Additive — Preserve M1 Values)

The M1 User entity already has a role ENUM with values `FREELANCER`, `CLIENT`, and `ADMIN`. M2 preserves these additively — do not rename or remove existing M1 role values.

Role	Permissions
CLIENT (M1 default)	Access all standard endpoints, manage own data, post jobs, review proposals
FREELANCER (M1)	Access all standard endpoints, manage own data, submit proposals, receive payouts
ADMIN (universal)	Everything CLIENT / FREELANCER can do, plus: change any user's role

- **Default role on registration:** CLIENT (Freelance's M1 default). The register payload may optionally include a role value of CLIENT or FREELANCER; if that field is missing, the server must default to CLIENT.
- ADMIN is never assigned on register
- Only an ADMIN can change another user's role via `PUT /api/users/{id}/role`
- The JWT token's role claim contains the user's actual role value
- Seed at least one ADMIN user so role management can be tested

5.4 Security Configuration (Per Service)

Each of the 5 services must have:

- A security configuration with a JWT filter that intercepts every request
- Stateless session management (no server-side sessions)
- CSRF disabled (since this is a token-based API)
- Public endpoints: `/api/auth/register` and `/api/auth/login` (on User Service only), and health checks
- All other endpoints require a valid JWT token

5.5 Password Handling

The User entity already has a password field from M1. In M2, you must ensure:

- Passwords are hashed with BCrypt before being stored
- The plaintext password is **never** stored or returned in API responses
- Login verifies the provided password against the stored hash

6 Database Architecture

6.1 Six Databases (Versions Match Task 2 Exactly)

All teams use the same 6 databases. **Docker image versions match Task 2** so students who completed Task 2 can reuse their stack without reconfiguration.

Database	Port	Docker Image	Architectural Role
PostgreSQL	5432	postgres:17	Primary relational data (from M1). PG 17 mandatory — PG 18 breaks Hibernate
MongoDB	27017	mongo:latest	Activity logs, event history, analytics — used by all services
Redis	6379	redis:latest	Caching layer for read-heavy endpoints — used by all services
Elasticsearch	9200	elasticsearch:8.19.12	Full-text search across jobs (S2). Pinned version.
Neo4j	7687	neo4j:latest	Recommendation graph from proposal patterns (S3)
Cassandra	9042	cassandra:latest	Time-series contract milestone events (S4)

6.2 Memory Limitations (Required)

Running 6 databases + 5 Spring Boot services strains laptop RAM. Apply these memory limits to avoid system freezes:

Service	Memory Limit	Rationale
Redis	256 MB max, allkeys-lru eviction	Bounded cache
Elasticsearch	512 MB JVM heap (ES_JAVA_OPTS=-Xms512m -Xmx512m)	Default heap is 4 GB
Cassandra	512 MB heap (MAX_HEAP_SIZE=512M, HEAP_NEWSIZE=128M)	Default is 1–2 GB
Neo4j	512 MB heap (NEO4J_server_memory_heap_max_size=512m)	Default can exceed 1 GB
PostgreSQL, MongoDB	Default	Well-behaved at defaults

Estimated total memory for the stack: PostgreSQL (~250 MB) + MongoDB (~400 MB) + Redis (256 MB capped) + Cassandra (~700 MB) + Neo4j (~600 MB) + Elasticsearch (~700 MB) + 5 Spring Boot apps ($5 \times 300 \text{ MB} = 1.5 \text{ GB}$) = **~4.5 GB** for the stack alone.

6.3 Dependency Model

Following the pattern from Lab 5 and Task 2:

- **PostgreSQL:** Hard dependency — the service will not start without it
- **MongoDB, Redis, Elasticsearch, Neo4j, Cassandra:** Soft dependencies — if any of these are unavailable, the features that depend on them should degrade gracefully (try-catch), but the service should still start and serve PostgreSQL-based endpoints

6.4 Docker Compose Additions

Add these containers to your existing docker-compose.yml alongside PostgreSQL and the 5 application services:

```
mongo:
  image: mongo:latest
  container_name: freelance-mongo
  ports:
    - "27017:27017"
  environment:
    MONGO_INITDB_ROOT_USERNAME: root
    MONGO_INITDB_ROOT_PASSWORD: rootpass
    MONGO_INITDB_DATABASE: freelancemongo
  volumes:
    - mongo-data:/data/db
  healthcheck:
    test: ["CMD", "mongosh", "--eval", "db.adminCommand('ping')"]
    interval: 10s
    timeout: 5s
    retries: 5

redis:
  image: redis:latest
  container_name: freelance-redis
  ports:
    - "6379:6379"
  command: redis-server --requirepass redispass --maxmemory 256mb --maxmemory-policy allkeys-lru
  healthcheck:
    test: ["CMD", "redis-cli", "-a", "redispass", "ping"]
    interval: 5s
    timeout: 5s
    retries: 5

elasticsearch:
  image: elasticsearch:8.19.12
  container_name: freelance-elasticsearch
  ports:
    - "9200:9200"
  environment:
    - discovery.type=single-node
    - xpack.security.enabled=false
    - ES_JAVA_OPTS=-Xms512m -Xmx512m
  volumes:
    - es-data:/usr/share/elasticsearch/data
  healthcheck:
    test: ["CMD-SHELL", "curl -f http://localhost:9200/_cluster/health || exit 1"]
    interval: 10s
    timeout: 5s
    retries: 10
    start_period: 30s

neo4j:
  image: neo4j:latest
  container_name: freelance-neo4j
  ports:
    - "7474:7474"
    - "7687:7687"
  environment:
    NEO4J_AUTH: neo4j/neo4jpass
    NEO4J_server_memory_heap_max_size: 512m
    NEO4J_server_memory_heap_initial_size: 256m
```



```

volumes:
  - neo4j-data:/data
healthcheck:
  test: ["CMD", "neo4j", "status"]
  interval: 10s
  timeout: 5s
  retries: 10
  start_period: 30s

cassandra:
  image: cassandra:latest
  container_name: freelance-cassandra
  ports:
    - "9042:9042"
  environment:
    CASSANDRA_CLUSTER_NAME: freelancecluster
    CASSANDRA_DC: datacenter1
    CASSANDRA_KEYSPACE: freelanceks
    MAX_HEAP_SIZE: 512M
    HEAP_NEWSIZE: 128M
  volumes:
    - cassandra-data:/var/lib/cassandra
  healthcheck:
    test: ["CMD-SHELL", "cqlsh -e 'DESCRIBE KEYSPACES'"]
    interval: 15s
    timeout: 10s
    retries: 10
    start_period: 60s

volumes:
  mongo-data:
  es-data:
  neo4j-data:
  cassandra-data:

```

6.5 Reference application.yml Fragments

Each of your 5 services needs the connection properties for every database it uses. Note: M1 used application.properties. In M2 you should migrate to application.yml (YAML format is what the auto-grader expects).

Shared across all services

```

spring:
  datasource:
    url: jdbc:postgresql://postgres:5432/freelancedb
    username: postgres
    password: postgres
  jpa:
    hibernate:
      ddl-auto: update
    show-sql: true
  data:
    redis:
      host: redis
      port: 6379
      password: redispass

```

```
jwt:
  secret: "<Base64-encoded shared secret across all 5 services>"
  expiration: 86400000 # 24 hours in ms
```

All services — MongoDB (event logging)

```
spring:
  data:
    mongodb:
      uri: mongodb://root:rootpass@mongo:27017/freelancemongo?authSource=admin
```

Job service only — Elasticsearch

```
spring:
  elasticsearch:
    uris: http://elasticsearch:9200
```

Proposal service only — Neo4j

```
spring:
  data:
    neo4j:
      uri: bolt://neo4j:7687
      username: neo4j
      password: neo4jpass
```

Contract service only — Cassandra

```
spring:
  cassandra:
    contact-points: cassandra
    port: 9042
    local-datacenter: datacenter1
    keyspace-name: freelanceks
    schema-action: CREATE_IF_NOT_EXISTS
```

7 New Entity Models

Your existing PostgreSQL entities from M1 remain unchanged except for the password handling described in Section 4.1 and Section 5.5. This section defines the **new** entities for the NoSQL databases.

7.1 MongoDB Document Entities (All Services)

Each service defines one MongoDB document class for its activity/event log. MongoDB entities must be **classes** (not records).

7.1.1 Common MongoEvent Interface

All 5 MongoDB event classes below implement a shared interface MongoEvent so that the EventFactory (Section 3.7) can return any concrete event typed as MongoEvent:

Method	Return type	Purpose
getId()	String	MongoDB ObjectId
getTimestamp()	LocalDateTime	When the event occurred
getAction()	String	The action identifier (e.g., REGISTERED)
getDetails()	Map<String, Object>	Additional event context

Each concrete event class below adds its own service-specific fields on top of this common interface.

7.1.2 AuthEvent (User Service)

Collection: auth_events

Field	Type	Constraints	Notes
id	String	auto-generated	MongoDB ObjectId
userId	Long	not null	References User in PG
action	String	not null	See action values below
timestamp	LocalDateTime	not null	When the event occurred
details	Map<String, Object>		Additional event context

action **primary values:** REGISTERED, LOGGED_IN, ROLE_CHANGED. **Non-exhaustive** — extend with domain-appropriate values when the Observer retrofit (Section 4.5) writes events from M1 user endpoints, e.g., USER_UPDATED (S1-F2 preferences), USER_DEACTIVATED (S1-F4), PRIMARY_SKILL_SET (S1-F7), USER_CREATED / USER_DELETED (CRUD). Use UPPER_SNAKE_CASE.

7.1.3 JobEvent (Job Service)

Collection: job_events

Field	Type	Constraints	Notes
id	String	auto-generated	MongoDB ObjectId
jobId	Long	not null	References Job in PG
action	String	not null	See action values below
timestamp	LocalDateTime	not null	
details	Map<String, Object>		Fields changed, dashboard params, etc.

action **primary values:** INDEXED, DASHBOARD_VIEWED. **Non-exhaustive** — extend for M1 retrofits, e.g., REQUIREMENTS_UPDATED (S2-F2), JOB_CLOSED (S2-F4), JOB_RATED (S2-F7), JOB_ATTACHMENT_VERIFIED (S2-F8), JOB_CREATED / JOB_DELETED (CRUD). Use UPPER_SNAKE_CASE.

7.1.4 ProposalEvent (Proposal Service)

Collection: proposal_events

Field	Type	Constraints	Notes
id	String	auto-generated	MongoDB ObjectId
proposalId	Long	not null	References Proposal in PG
action	String	not null	See action values below
timestamp	LocalDateTime	not null	
details	Map<String, Object>		

action **primary values:** ANALYTICS_VIEWED, INTERACTION_RECORDED. **Non-exhaustive** — extend for

M1 retrofits, e.g., PROPOSAL_ACCEPTED (S3-F2), PROPOSAL_COMPLETED (S3-F4), PROPOSAL_WITHDRAWN (S3-F7), MILESTONES_ADDED (S3-F8), PROPOSAL_CREATED / PROPOSAL_DELETED (CRUD). Use UPPER_SNAKE_CASE.

7.1.5 ContractEvent (Contract Service)

Collection: contract_events

Field	Type	Constraints	Notes
id	String	auto-generated	MongoDB ObjectId
contractId	Long	not null	References Contract in PG
action	String	not null	See action values below
timestamp	LocalDateTime	not null	
details	Map<String, Object>		

action **primary values:** MILESTONE_TRACKED, ANALYTICS_VIEWED. **Non-exhaustive** — extend for M1 retrofits, e.g., PROGRESS_UPDATED (S4-F2), BATCH_STATUS_UPDATED (S4-F4), OLD_DATA_PURGED (S4-F7), CONTRACT_DELETED (CRUD). Use UPPER_SNAKE_CASE.

7.1.6 PayoutAuditEvent (Wallet Service)

Collection: payout_audit_trail

Field	Type	Constraints	Notes
id	String	auto-generated	MongoDB ObjectId
payoutId	Long	not null	References Payout in PG
action	String	not null	See action values below
timestamp	LocalDateTime	not null	
method	String	conditional	See method/amount note below
amount	Double	conditional	See method/amount note below
details	Map<String, Object>		Reversal reason, strategy name, etc.

action **primary values:** CREATED, COMPLETED, FAILED, REFUNDED, REFUND_DENIED, ANALYTICS_VIEWED. **Non-exhaustive** — extend for M1 retrofits, e.g., PROMO_APPLIED (S5-F5), RETRY_ATTEMPTED (S5-F7), PAYOUT_DELETED (CRUD). Use UPPER_SNAKE_CASE.

method **values:** BANK_TRANSFER, PAYPAL, CRYPTO — matches the M1 Payout.method enum values.

When each action value is written:

- CREATED — written by M1 S5-F4 (Process Payout) when the Payout row is first inserted.
- COMPLETED — written by M1 S5-F4 when the payout status transitions to COMPLETED.
- FAILED — written by M1 S5-F4 when a simulated gateway failure occurs. Since M1 does not model a real payment gateway, enable failure simulation via an optional query parameter ?simulateFailure=true on the S5-F4 create-payout request: when set, S5-F4 short-circuits to Payout.status = FAILED in PG and writes a FAILED audit event. This affordance exists solely to let the S5-F11 test scenario produce failed-payout data; it is not a production flow.
- REFUNDED — written by M1 S5-F2 (simple refund) and by M2 S5-F12 (milestone-based reversal) when the reversal is successfully processed.
- REFUND_DENIED — written by M2 S5-F12 when NoReversalStrategy blocks the reversal (e.g., the reversal window has expired). The event records the denial reason and strategy name in details.
- ANALYTICS_VIEWED — written by M2 S5-F10 (Platform Fee Analytics) each time the endpoint is invoked, regardless of whether the response is served from cache. Used to track analytics access patterns.

method **and** amount **field requirement:** The fields method and amount are **required (not null)** on every event whose action is a payout-shaped action, i.e., $\text{action} \in \{\text{CREATED, COMPLETED, FAILED, REFUNDED, REFUND_DENIED}\}$ and any payout-lifecycle retrofit action (e.g., `PROMO\_APPLIED`, `RETRY\_ATTEMPTED`). They are **omitted (null-permitted)** only on non-payout actions such as `ANALYTICS\_VIEWED` (an analytics view is not tied to a specific Payout row's method/amount). S5-F11 (Payout Method Breakdown) relies on this: it groups and sums only events whose method and amount are populated, so a `CREATED/COMPLETED/FAILED` event written without method would silently vanish from the breakdown.

7.2 Elasticsearch Document (Job Service)

7.2.1 JobSearchDocument

Index: jobs

Field	ES Field Type	Purpose
id	Keyword	Exact match, corresponds to PG id
title	Text (analyzed)	Full-text search on job title
description	Text (analyzed)	Full-text search on job description
category	Keyword	Filtering (<code>WEB_DEV</code> , <code>MOBILE</code> , <code>DESIGN</code> , <code>WRITING</code>)
budgetMin	Double	Range filtering
budgetMax	Double	Range filtering
rating	Double	Range filtering
status	Keyword	Filtering (<code>OPEN</code> , <code>IN_PROGRESS</code> , <code>CLOSED</code> — matches M1 Job.status enum)

Note on title and description: Both fields are **top-level columns** on the M1 Job entity (not JSONB keys). S2-F11 reads them directly from the PG jobs table without any JSONB lookup — no additive M1 schema change is needed for the search document.

7.3 Neo4j Entities (Proposal Service)

7.3.1 FreelancerNode

Property	Type	Notes
userId	Long	Corresponds to User.id in PG (users with role <code>FREELANCER</code>)
name	String	Freelancer's display name

7.3.2 JobNode

Property	Type	Notes
jobId	Long	Corresponds to Job.id in PG
title	String	Job title
category	String	For filtering recommendations

7.3.3 PROPOSED_TO Relationship

Connects a FreelancerNode to a JobNode. Direction: (Freelancer)-[:PROPOSED_TO]->(Job).

Property	Type	Notes
proposalCount	Integer	Incremented each time freelancer submits a proposal on this job
lastProposalDate	LocalDateTime	Updated on each new proposal

7.4 Cassandra Entity (Contract Service)

7.4.1 ContractMilestoneEvent

Table: contract_milestone_events

Column	Type	Key	Notes
contract_id	Long	Partition Key	Groups all milestone events for one contract
timestamp	Timestamp	Clustering (DESC)	Orders events newest-first
milestone_order	Integer		1, 2, 3... — matches ProposalMilestone.milestoneOrder from M1
status	String		PENDING, IN_PROGRESS, COMPLETED, APPROVED (matches M1 ProposalMilestone.status)
recorded_by	String		User who recorded the event (freelancer or client name)
notes	String		Optional event notes

Queries on this table **must** include contract_id in the WHERE clause (partition key requirement).

8 Caching Strategy

All read-heavy endpoints (both new M2 features and existing M1 endpoints) must be cached in Redis.

8.1 TTL Guidelines

Data Type	TTL	Examples
Dashboards / analytics	10 minutes	Revenue dashboard, performance dashboard
Search results	5 minutes	Job search, recommendations
Activity feeds	5 minutes	User activity feed
Entity detail views	15 minutes	Job profile, contract details

8.2 Invalidation

Write operations (POST, PUT, DELETE) must invalidate any cached data they affect. The canonical invalidation enumeration for all M1 write endpoints + CRUD writes + M2 observer writes lives in Section 4.4.4 and uses the wildcard-deletion strategy described in Section 4.4.6. Features that mention invalidation (e.g., S5-F12 step h) refer to those rules. In brief:

- Entity write on <entity>{id} → delete <service>::<entity>::{id} plus wildcard-delete <service>::S{n}-F{m}::* for every feature whose cached output may reference that entity.
- M2 analytics events written by Observer retrofit (see Section 4.4.4 M2-writes subsection) also invalidate the matching dashboard caches.
- Over-invalidation (clearing a few cache keys that did not strictly need to be cleared) is acceptable — correctness beats cache-hit optimality at M2.

9 Cross-Cutting Requirements

These requirements apply globally and are **not** counted as numbered features. Each requirement has a dedicated test scenario below.

9.1 CC-1 — JWT on All Endpoints

Requirement: Every endpoint (including your M1 endpoints) must require a valid JWT token in the Authorization header, except POST /api/auth/register, POST /api/auth/login, and health checks. Missing or invalid token returns **401**. Insufficient role returns **403**.

Test scenario:

- a) Enumerate every endpoint in all 5 services by scanning controllers. For each: call without Authorization header → expect 401 unless the endpoint is /api/auth/register, /api/auth/login, or a health check.
- b) For each public endpoint: call without a token → expect a 2xx response (not 401).
- c) Call any protected endpoint with a malformed token (Bearer abc) → 401.
- d) Call any protected endpoint with an expired token → 401.
- e) Call PUT /api/users/{id}/role with a valid CLIENT token → 403.
- f) Count the protected vs public endpoints. Freelance should have exactly 3 public endpoints (register, login, health). Anything else public is a grading failure.

9.2 CC-2 — Role Management

Requirement: Implement PUT /api/users/{id}/role as an ADMIN-only endpoint that changes a user's role. Must log a ROLE_CHANGED event to auth_events.

Request body: {"role": "ADMIN"} (or "CLIENT" or "FREELANCER").

Behavior:

- a) Validate JWT and role claim is ADMIN — **throws 403 if not ADMIN**.
- b) Find user by ID — **throws 404 if user not found**.
- c) Validate requested role is a valid ENUM value (CLIENT, FREELANCER, or ADMIN) — **throws 400 if invalid**.
- d) Update the user's role and save.
- e) Log ROLE_CHANGED to auth_events with the old and new role in details.
- f) Invalidate cached user detail. **Note:** The ROLE_CHANGED write to auth_events automatically triggers §4.4.4 invalidation of user-service::S1-F12:* (the activity feed) via the Observer chain; no extra code is needed in this endpoint.
- g) Return the updated user with status **code 200**.

Token staleness after role change (accepted limitation): The target user's existing JWT continues to carry its **old** role claim until it expires. Since tokens have a 24-hour lifetime (see §5.2) and the spec does not require a token-revocation list, a demoted user (ADMIN → CLIENT / FREELANCER) retains ADMIN authorization for up to 24 hours, and a promoted user (CLIENT / FREELANCER → ADMIN) cannot exercise new ADMIN privileges until they log in again. This is an accepted limitation for M2. The ROLE_CHANGED audit event and the invalidated user-detail cache give operators the information they need, but do not revoke the in-flight JWT. Production-grade solutions (token revocation list, shorter expiry + refresh tokens, per-request PG role re-check) are out of scope for M2.

Test scenario:

- a) Seed an ADMIN and a CLIENT. Log in as ADMIN.
- b) PUT /api/users/{clientId}/role with body {"role":"ADMIN"} → 200, the target user's role is now ADMIN in PG.
- c) Verify the auth_events collection has a ROLE_CHANGED event with details containing the old and new roles.
- d) Verify the cached user detail (user-service::user::{clientId}) was removed from Redis.

- e) Verify cached keys under `user-service::S1-F12::{clientId}::*` were removed from Redis after the `ROLE_CHANGED` write propagated through the Observer chain (Â§4.4.4 NoSQL-writer invalidation).
- f) Log in as a different CLIENT and call `PUT /api/users/1/role` → 403.
- g) Call `PUT /api/users/999/role` with ADMIN token → 404.
- h) Call with body `{"role":"BANANA"}` → 400.
- i) Call without Authorization → 401.

9.3 CC-3 — Redis Caching on M1 Endpoints

Requirement: Your M1 feature GET endpoints (27 for Freelance — see Section 4.4.1) and CRUD GET `/api/<entity>/<id>` endpoints (10 for Freelance) must be cached per Section 4.4. List endpoints are **not** cached.

Test scenario: See Section 4.4.6 for the comprehensive cache test scenario. Additional spot checks:

- a) For each of the 27 cached M1 feature endpoints and 10 CRUD GET-by-ID endpoints: call twice, inspect Redis between calls, confirm a key with the expected `<service>::<featureId>::<param-hash>` or `<service>::<entity>::<id>` format exists.
- b) For every M1 write endpoint (the 18 feature writes + 30 CRUD writes): trigger the write, then query Redis and verify the affected detail cache key is gone and relevant feature-result caches are invalidated.
- c) Confirm GET `/api/<entity>` (list) endpoints produce NO cache entry.

9.4 CC-4 — Design Pattern Implementation

Requirement: The 7 design patterns defined in Section 3 must be implemented at their specified locations.

Test scenario: See each pattern's own test scenario in Section 3.2 through 3.8.

9.5 CC-5 — Docker Compose with 6 Databases

Requirement: Your `docker-compose.yaml` must include all 6 database containers (PostgreSQL, MongoDB, Redis, Elasticsearch, Neo4j, Cassandra) with versions matching Task 2 alongside your 5 application services.

Test scenario:

- a) Parse the team's `docker-compose.yaml`. Assert services named `postgres`, `mongo`, `redis`, `elasticsearch`, `neo4j`, `cassandra` exist.
- b) Assert image tags: `postgres:17`, `elasticsearch:8.19.12`, and `mongo:latest/redis:latest/neo4j:latest/cassandra:latest`.
- c) Assert memory caps: Redis command includes `-maxmemory 256mb -maxmemory-policy allkeys-lru`; Elasticsearch env has `ES_JAVA_OPTS=-Xms512m -Xmx512m`; Cassandra env has `MAX_HEAP_SIZE: 512M`; Neo4j env has `NEO4J_server_memory_heap_max_size: 512m`.
- d) Assert healthchecks exist for all databases.
- e) Run `docker compose up -d` → all 6 databases reach healthy state within 120 seconds.
- f) All 5 application services start successfully and connect to their databases.
- g) Total stack memory usage (`docker stats`) stays under 5 GB.

9.6 CC-6 — Application Configuration (application.yml)

Requirement: Each service's application.yml must include connection settings for the databases it uses, plus the shared JWT secret and expiration. Must be YAML format, not properties.

Test scenario:

- a) Assert every service has an application.yml file (not application.properties).
- b) For user-service, job-service, proposal-service, contract-service, wallet-service: assert spring.datasource.url is present and points to postgres:5432.
- c) Assert spring.data.mongodb.uri, spring.data.redis.host, and jwt.secret/jwt.expiration are present in every service.
- d) Job-service: assert spring.elasticsearch.uris is present.
- e) Proposal-service: assert spring.data.neo4j.uri is present.
- f) Contract-service: assert spring.cassandra.contact-points and keyspace-name are present.
- g) Boot each service in isolation with other services down → must start as long as PostgreSQL is up (hard dep). NoSQL unavailability must not prevent startup.

10 Features

Milestone 2 adds **15 new features** (3 per service). Feature numbering continues from M1 (which used F1–F9), so M2 features are **F10–F12** per service.

Every M2 feature touches at least 2 databases. Each feature specification lists which databases are involved, whether authentication is required, and the caching behavior.

Important Note: Do not forget to apply the Layered architecture (having repository, service and controller) as the test case will test that each feature is implemented according to the layered architecture as explained to you in the Labs.

10.1 User Service Features (port 8081)

10.1.1 [S1-F10] Register User

Branch: feat/user/S1-F10/<ID>

Endpoint: POST /api/auth/register

Auth: None (public endpoint)

Databases: PostgreSQL, MongoDB

Request body:

```
{
  "name": "Ahmed Ali",
  "email": "ahmed@example.com",
  "password": "securePassword123",
  "phone": "+201012345678",
  "role": "FREELANCER"
}
```

Response:

```
{
  "token": "eyJhbGciOiJIUzI1NiJ9...",
  "expiresIn": 86400000
}
```

Behavior:

- a) Validate that name, email, password, and phone are not blank –**throws 400 if any field is missing or blank**.
- b) Check that no user with this email or phone already exists –**throws 409 if the email or phone is already registered** (the M1 User entity requires both email and phone to be unique).
- c) Hash the password before saving with **BCrypt** (the stored password must not be the plaintext value).
- d) Determine the role to assign: if the request body contains a role value of CLIENT or FREELANCER, honour it; if the body contains role=ADMIN or omits role, default to CLIENT (Freelance's M1 default role — never assign ADMIN on register).
- e) Create the user in PostgreSQL with the resolved role and status ACTIVE.
- f) Log a REGISTERED event to the auth_events collection in MongoDB with the user's ID and the current timestamp.
- g) Generate a token containing the user's email (as the sub claim), the user's numeric User.id (as a custom uid claim), and the user's role (as a role claim). The uid claim is required by the ownership checks in Å§10 (S1-F12, S3-F12); see Å§5.2 for the full payload contract.
- h) Return the token and expiration with status **code 201**.

Test scenario:

- a) POST /api/auth/register with valid data (no role field) → 201 with token. Verify user exists in database with role CLIENT.
- b) POST /api/auth/register with "role":"FREELANCER" → 201; the new row's role is FREELANCER.
- c) POST /api/auth/register with "role":"ADMIN" → 201 but the new row's role is CLIENT (ADMIN silently ignored).
- d) POST /api/auth/register with the same email → 409.
- e) POST /api/auth/register with the same phone (but different email) → 409.
- f) POST /api/auth/register with blank email or blank phone → 400.
- g) Verify the stored password is **not** the plaintext password (it should be a BCrypt hash).

10.1.2 [S1-F11] Login

Branch: feat/user/S1-F11/<ID>

Endpoint: POST /api/auth/login

Auth: None (public endpoint)

Databases: PostgreSQL, MongoDB

Request body:

```
{
  "email": "ahmed@example.com",
  "password": "securePassword123"
}
```

Response:

```
{
  "token": "eyJhbGciOiJIUzI1NiJ9...",
  "expiresIn": 86400000
}
```

Behavior:

- a) Find user by email in PostgreSQL –**throws 401 if user not found**.
- b) Verify the provided password against the stored hash –**throws 401 if the password does not match**.
- c) Log a LOGGED_IN event to the auth_events collection in MongoDB with the user's ID and timestamp.
- d) Generate a token containing the user's email (as the sub claim), the user's numeric User.id (as a custom uid claim), and the user's role (as a role claim). The uid claim is required by the ownership checks in §10 (S1-F12, S3-F12); see §5.2 for the full payload contract.
- e) Return the token and expiration with status **code 200**.

Note on 401 for both error cases: Returning **401** for both “user not found” and “wrong password” is intentional — it prevents account enumeration. A caller cannot distinguish whether a given email is registered. Do **not** return 404 for missing email.

Test scenario:

- a) Register a user, then POST /api/auth/login with correct credentials → 200 with token.
- b) POST /api/auth/login with wrong password → 401.
- c) POST /api/auth/login with non-existent email → 401 (not 404 — this is deliberate to prevent account enumeration).
- d) Use the returned token to access a protected endpoint (e.g., GET /api/users/1) → should succeed. Access the same endpoint without a token → 401.

10.1.3 [S1-F12] Get User Activity Feed

Branch: feat/user/S1-F12/<ID>

Endpoint: GET /api/users/{id}/activity?page={page}&size={size}

Auth: Required (USER)

Databases: MongoDB, PostgreSQL, Redis

Response: Paginated list of activity events:

```
{
  "content": [
    {
      "action": "LOGGED_IN",
      "timestamp": "2026-04-20T14:30:00",
      "details": {}
    },
    {
      "action": "REGISTERED",
      "timestamp": "2026-04-20T14:25:00",
      "details": {"email": "ahmed@example.com"}
    }
  ],
  "page": 0,
  "size": 10,
  "totalElements": 2
}
```

Behavior:

- a) Validate the JWT token –**throws 401 if missing or invalid**.

- b) **Ownership check:** Verify the authenticated caller's uid claim from the JWT equals the path {id}, OR the caller's role claim is ADMIN –**throws 403 if the caller is neither the target user nor an ADMIN** (prevents account enumeration and PII leakage). The comparison is a direct numeric equality check on the JWT's uid claim — no PG lookup needed.
- c) Find user by ID in PostgreSQL –**throws 404 if user not found**.
- d) Query the auth_events collection in MongoDB for events belonging to this user, sorted by most recent first, paginated by the page and size parameters. **Defaults and bounds:** if page is omitted, default to 0; if size is omitted, default to 10; cap size at 100 — requests with size > 100 are silently clamped to 100.
- e) The response should be cached for 5 minutes. Subsequent identical requests within the cache window should return the cached version.
- f) Return the paginated activity feed with status **code 200**.

Note: The activity feed should also include events from role management operations. The role management endpoint (PUT /api/users/{id}/role, see Section 9 cross-cutting requirement CC-2) must also be implemented, and ROLE_CHANGED events should appear in this feed.

Test scenario:

- a) Register a user, login, then GET /api/users/{id}/activity with the user's own token → should contain REGISTERED and LOGGED_IN events ordered by most recent first.
- b) Using the seeded ADMIN token, trigger a role change via PUT /api/users/{id}/role changing the user's role from CLIENT to FREELANCER. Then GET /api/users/{id}/activity with the affected user's token → should contain a ROLE_CHANGED event as the most recent entry.
- c) Register user A and user B. With user A's token, call GET /api/users/{B.id}/activity → 403 (ownership violation, not 404 — A's token is valid and B exists).
- d) With the seeded ADMIN token, call GET /api/users/{any-user-id}/activity → 200 (ADMIN bypasses ownership check).
- e) GET /api/users/{id}/activity without Authorization header → 401.
- f) GET /api/users/999/activity with an ADMIN token → 404 (ADMIN passes ownership, user not found).
- g) Verify pagination: create multiple events, request with page=0&size=1 (with matching-owner token) → should return 1 event with correct total count.

10.2 Job Service Features (port 8082)

10.2.1 [S2-F10] Full-Text Job Search

Branch: feat/job/S2-F10/<ID>

Endpoint:

GET /api/jobs/search/full-text?

query={text}&category={cat}&status={s}&minBudget={n}&maxBudget={n}

Auth: Required (USER)

Databases: Elasticsearch, PostgreSQL, Redis

Response: List of matching jobs with relevance-based ordering.

Note on path: This is a **distinct endpoint** from M1's S2-F1 (GET /api/jobs/search). M1's endpoint searches PostgreSQL by status and budget range and is unchanged; M2's new /full-text endpoint searches Elasticsearch on title and description. Both endpoints must coexist.

Behavior:

- a) Validate the JWT token *–throws 401 if missing or invalid.*
- b) Search jobs using full-text search on the query parameter, matching against the job's title and description. The search should handle partial matches and be case-insensitive.
- c) Optionally filter by category (exact match, one of WEB_DEV, MOBILE, DESIGN, WRITING), status (exact match, one of OPEN, IN_PROGRESS, CLOSED), and minBudget/maxBudget (range filter against budgetMin and budgetMax). All filter parameters are optional — if not provided, they are ignored.
- d) Results should be sorted by relevance to the search query.
- e) The response should be cached for 5 minutes.
- f) Return the matching jobs with status **code 200**. Return an empty list if no matches.

Test scenario:

- a) Create 3 jobs via CRUD and index them. Search with query=react → should return jobs with “react” in title or description.
- b) Search with query=react&category=WEB_DEV → should return only WEB_DEV react jobs.
- c) Search with query=react&status=OPEN → should return only OPEN react jobs.
- d) Search with query=react&minBudget=1000 → should return only jobs with budget range overlapping 1000+.
- e) Search with query=nonexistent → empty list.
- f) Search without token → 401.

10.2.2 [S2-F11] Index Job for Search

Branch: feat/job/S2-F11/<ID>

Endpoint: POST /api/jobs/{id}/index

Auth: Required (USER)

Databases: PostgreSQL, Elasticsearch

Behavior:

- a) Validate the JWT token *–throws 401 if missing or invalid.*
- b) Find job by ID in PostgreSQL *–throws 404 if job not found.*
- c) Create or update the job's search document in Elasticsearch with the job's id, title, description, category, budgetMin, budgetMax, rating, and status. All fields are read directly from the PG jobs top-level columns — no JSONB lookup is needed (title and description are top-level columns on the M1 Job entity).
- d) Log an INDEXED event to the job_events collection in MongoDB via the Observer chain. The event's details object must include jobId, the list of indexedFields written in the prior step, and a source key: "explicit" when invoked through this endpoint, "auto_crud_create" / "auto_crud_update" when triggered by the CRUD auto-index retrofit. On DELETE the retrofit emits JOB_DELETED (not INDEXED) after removing the ES document.
- e) **Auto-index on CRUD changes:** In addition to this explicit endpoint, the job must be automatically re-indexed whenever it is created or updated via any M1 CRUD endpoint, and **removed from the search index** whenever it is deleted via DELETE /api/jobs/{id}. This is graded — tests will create/update a job via CRUD and then search for it without calling this endpoint explicitly; after a delete, searching for the deleted job must return no match.
- f) Return status **code 200** indicating successful indexing.

Test scenario:

- a) Create a job via CRUD with title="React Native app" and description="build a cross-platform mobile app", then POST /api/jobs/{id}/index → 200. Search with query=react → should return the job.
- b) POST /api/jobs/999/index → 404.
- c) Update a job's title via CRUD (without calling /index), then search with the new title → should find the updated job (auto-indexing verification).
- d) Delete a job via CRUD, then search for it → should return no match (auto-remove verification).
- e) Index without token → 401.

10.2.3 [S2-F12] Get Job Market Dashboard

Branch: feat/job/S2-F12/<ID>

Endpoint: GET /api/jobs/{id}/dashboard

Auth: Required (USER)

Databases: PostgreSQL, MongoDB, Redis

Response DTO (JobDashboardDTO): jobId, title, totalProposals, acceptedProposals, averageBidAmount, activeAttachments, rating.

Behavior:

- a) Validate the JWT token **–throws 401 if missing or invalid.**
- b) Find job by ID in PostgreSQL **–throws 404 if job not found.**
- c) Aggregate from the proposals table: total number of proposals for this job, number of accepted proposals (status = ACCEPTED), and average bid amount.
- d) Count active attachments for this job from the job_attachments table where expiryDate >= CURRENT_DATE.
- e) Include the job's current rating from the jobs table.
- f) Log a DASHBOARD_VIEWED event to the job_events collection in MongoDB. **This log must be written on every invocation**, independently of whether the response was served from cache — perform the logging step **outside** the cache decorator/layer so it runs on cache hits too. Note that ANALYTICS_VIEWED / DASHBOARD_VIEWED are pure-observability actions and are excluded from Observer-driven cache invalidation (see §4.4.4) — they do not mutate the data that this dashboard reads, so invalidating the cache on them would defeat the cache.
- g) The response should be cached for 10 minutes.
- h) Return the dashboard DTO with status **code 200**.

Test scenario:

- a) Create a job with 2 attachments and 5 proposals (bids 500, 800, 1000, 1200, 900; 1 ACCEPTED). GET /api/jobs/{id}/dashboard → totalProposals=5, acceptedProposals=1, averageBidAmount=880, activeAttachments=2.
- b) GET /api/jobs/999/dashboard → 404.
- c) Job with no proposals → totalProposals=0, acceptedProposals=0, averageBidAmount=0.
- d) Without token → 401.

10.3 Proposal Service Features (port 8083)

10.3.1 [S3-F10] Get Proposal Analytics Dashboard

Branch: feat/proposal/S3-F10/<ID>

Endpoint: GET /api/proposals/analytics/dashboard?startDate={date}&endDate={date}

Auth: Required (USER)

Databases: PostgreSQL, MongoDB, Redis

Response DTO (ProposalAnalyticsDashboardDTO): totalProposals, acceptanceRate, averageBidAmount, averageEstimatedDays, proposalsByStatus (map of status to count).

Note on path and DTO: This is a **distinct endpoint and DTO** from M1's S3-F6 (GET /api/proposals/analytics). M1's endpoint returns a summary DTO and remains unchanged; M2's new /dashboard endpoint returns the richer ProposalAnalyticsDashboardDTO with a status breakdown map. Both must coexist — do not overwrite either.

Behavior:

- Validate the JWT token —**throws 401 if missing or invalid**.
- Validate date range —**throws 400 if startDate is after endDate**. Both startDate and endDate are LocalDate values on the wire; internally expand the filter window to the fully-closed range [startDate T00:00:00, endDate T23:59:59.999] in the server time zone, matching the PG TIMESTAMP columns (submittedAt, createdAt, etc.) and the MongoDB timestamp field. Boundary rows (events exactly at startDate T00:00:00 or endDate T23:59:59.999) are included.
- Aggregate proposal statistics from the proposals table for the given date range (filter on submittedAt): total proposals, average bid amount, average estimated days, acceptance rate (accepted proposals divided by total proposals), and a breakdown of proposal counts by status.
- Log an ANALYTICS_VIEWED event to the proposal_events collection in MongoDB. **This log must be written on every invocation**, independently of whether the response was served from cache — perform the logging step **outside** the cache decorator/layer so it runs on cache hits too. Note that ANALYTICS_VIEWED / DASHBOARD_VIEWED are pure-observability actions and are excluded from Observer-driven cache invalidation (see §4.4.4) — they do not mutate the data that this dashboard reads, so invalidating the cache on them would defeat the cache.
- The response should be cached for 10 minutes.
- Return the analytics DTO with status **code 200**.

Test scenario:

- Create 10 proposals in March 2026: 4 ACCEPTED, 2 REJECTED, 2 WITHDRAWN, 2 SUBMITTED. GET /api/proposals/analytics/dashboard?startDate=2026-03-01&endDate=2026-03-31 → totalProposals=10, acceptanceRate=0.4, proposalsByStatus={ACCEPTED:4, REJECTED:2, WITHDRAWN:2, SUBMITTED:2}.
- Query with dates where no proposals exist → all values 0.
- startDate=2026-04-01&endDate=2026-03-01 → 400.
- Without token → 401.

10.3.2 [S3-F11] Record Freelancer-Job Interaction

Branch: feat/proposal/S3-F11/<ID>

Endpoint: POST /api/proposals/{proposalId}/record-interaction

Auth: Required (USER)

Databases: Neo4j, PostgreSQL, MongoDB

Behavior:

- a) Validate the JWT token –**throws 401 if missing or invalid**.
- b) Find the proposal by ID in PostgreSQL –**throws 404 if proposal not found**.
- c) Verify the proposal's status is SUBMITTED –**throws 400 if the proposal is not in SUBMITTED status** (meaning we only record interactions for live proposals).
- d) **Idempotency check:** Before mutating the graph, check whether this specific proposalId has already been recorded. Maintain the idempotency marker **inside Neo4j** (not PostgreSQL – M2 does not alter any M1 PG schema per Å§1). Two acceptable Neo4j-only approaches: (a) a recorded_proposal_ids collection property on the PROPOSED_TO relationship that accumulates the set of proposalIds contributing to the edge's proposalCount, or (b) a dedicated Neo4j idempotency sentinel node, e.g., (Freelancer)-[:RECORDED_PROPOSAL {proposalId:...}]->(Job), checked with an EXISTS query. If the marker says this proposal was already recorded, **return 200 immediately without mutating proposalCount or lastProposalDate**. This guarantees the endpoint is idempotent under repeat POSTs without introducing any M1 PostgreSQL schema change.
- e) Look up the freelancer who submitted this proposal and the job the proposal targets using the M1 cross-service native SQL pattern – query the users and jobs tables directly via the shared PostgreSQL database using the proposal's foreign key columns (freelancer_id, job_id). **Do not introduce inter-service HTTP calls**; those come in Milestone 3.
- f) In the recommendation graph: find or create a node for the freelancer and a node for the job. Then find or create a PROPOSED_TO relationship between them. If the relationship already exists, increment the proposalCount by 1 and update the lastProposalDate. If it is new, set proposalCount to 1. Record this proposalId in the idempotency marker from step d so future calls with the same proposalId are no-ops.
- g) Log an INTERACTION_RECORDED event to the proposal_events collection in MongoDB via the Observer chain. The event's details object must include proposalId, freelancerId, and jobId. This write is what causes the S3-F12 recommendation cache to be wildcard-invalidated per Å§4.4.4 (NoSQL-writer → cached-reader rule); it is therefore required on the non-idempotent path only – on the idempotent short-circuit at step d, the graph did not change, so do **not** emit this event.
- h) Return status **code 200** with a confirmation message.

Test scenario:

- a) Create a freelancer, a job, and a SUBMITTED proposal P1. POST /api/proposals/P1/record-interaction → 200, proposalCount=1.
- b) Call the same endpoint again on the same P1 → 200, but proposalCount is still 1 (idempotency – repeated calls on the same proposalId must not inflate the counter).
- c) Create a second SUBMITTED proposal P2 from the same freelancer on the same job. POST /api/proposals/P2/record-interaction → 200, proposalCount=2.
- d) Call with a WITHDRAWN (not submitted) proposal → 400.
- e) Call with non-existent proposal ID → 404.
- f) Without token → 401.

10.3.3 [S3-F12] Get Recommended Jobs for Freelancer

Branch: feat/proposal/S3-F12/<ID>

Endpoint: GET /api/proposals/recommendations?freelancerId={id}&limit={n}

Auth: Required (USER)

Databases: Neo4j, PostgreSQL, Redis

Response: List of JobRecommendationDTO: jobId, title, category, score (how many similar freelancers proposed to this job).

Behavior:

- a) Validate the JWT token *–throws 401 if missing or invalid.*
- b) **Ownership check:** Verify the authenticated caller's uid claim from the JWT equals the freelancerId query parameter, OR the caller's role claim is ADMIN *–throws 403 if the caller is neither the target freelancer nor an ADMIN* (recommendations reveal proposal-pattern adjacency — exposing them to arbitrary callers leaks behavioral data). The comparison is a direct numeric equality check on the JWT's uid claim — no PG lookup needed.
- c) Find freelancer by ID in PostgreSQL *–throws 404 if user not found.*
- d) Traverse the recommendation graph: find other freelancers who proposed to the same jobs as this freelancer (meaning they share at least one job in common), then find jobs those other freelancers proposed to that this freelancer has **not** proposed to yet. Rank the results by how many of those similar freelancers proposed to each recommended job.
- e) Enrich the results with job details (title, category) from PostgreSQL.
- f) Limit results to the top limit recommendations. Default limit to 5 if not provided.
- g) The response should be cached for 5 minutes.
- h) Return the recommendations with status **code 200**. Return an empty list if no recommendations can be made.

Test scenario:

- a) Create 3 freelancers (A, B, C) and 4 jobs (J1, J2, J3, J4). Record interactions: A→J1, A→J2; B→J1, B→J3; C→J2, C→J4. Get recommendations for A with A's own token → should include J3 (because B also proposed to J1) and J4 (because C also proposed to J2). Should NOT include J1 or J2.
- b) Get recommendations for A with B's token → 403 (ownership violation).
- c) Get recommendations for A with an ADMIN token → 200 (ADMIN bypass).
- d) Get recommendations for a freelancer with no recorded interactions (own token) → empty list.
- e) freelancerId=999 with ADMIN token → 404.
- f) Without token → 401.

10.4 Contract Service Features (port 8084)

10.4.1 [S4-F10] Get Contract Analytics Dashboard

Branch: feat/contract/S4-F10/<ID>

Endpoint: GET /api/contracts/analytics?startDate={date}&endDate={date}

Auth: Required (USER)

Databases: PostgreSQL, MongoDB, Redis

Response DTO (ContractAnalyticsDTO): totalContracts, averageContractValue, completionRate, averageContractDurationDays, contractsByStatus (map of status to count).

Behavior:

- a) Validate the JWT token *–throws 401 if missing or invalid.*

- b) Validate date range –**throws 400 if startDate is after endDate**. Both startDate and endDate are LocalDate values on the wire; internally expand the filter window to the fully-closed range [startDate T00:00:00, endDate T23:59:59.999] in the server time zone, matching the PG TIMESTAMP columns (startDate, createdAt, etc.) and the MongoDB timestamp field. Boundary rows (events exactly at startDate T00:00:00 or endDate T23:59:59.999) are included.
- c) Aggregate from the contracts table for the given date range (filter on startDate).
- d) **totalContracts**: count of contracts whose startDate is within [startDate, endDate].
- e) **averageContractValue**: average of agreedAmount across all contracts in the range. Return 0 if none.
- f) **completionRate**: COMPLETED contracts / totalContracts in the range. Return 0 if total is 0.
- g) **averageContractDurationDays**: for contracts with status = COMPLETED and non-null endDate, compute the average of (endDate - startDate) in days. Return 0 if none completed in the range.
- h) **contractsByStatus**: count of contracts grouped by their current status column (ACTIVE, COMPLETED, TERMINATED, DISPUTED).
- i) Log an ANALYTICS_VIEWED event to the contract_events collection in MongoDB. **This log must be written on every invocation**, independently of whether the response was served from cache – perform the logging step **outside** the cache decorator/layer so it runs on cache hits too. Note that ANALYTICS_VIEWED / DASHBOARD_VIEWED are pure-observability actions and are excluded from Observer-driven cache invalidation (see §4.4.4) – they do not mutate the data that this dashboard reads, so invalidating the cache on them would defeat the cache.
- j) The response should be cached for 10 minutes.
- k) Return the analytics DTO with status **code 200**.

Test scenario:

- a) Create 8 contracts in April 2026: 5 COMPLETED (with endDate 30, 45, 60, 30, 45 days after startDate), 2 ACTIVE, 1 TERMINATED. GET /api/contracts/analytics?startDate=2026-04-01&endDate=2026-04-30 → totalContracts=8, completionRate=0.625 (5 of 8), averageContractDurationDays=42.
- b) Dates with no contracts → totalContracts=0, averageContractValue=0, completionRate=0, empty contractsByStatus.
- c) Invalid date range → 400.
- d) Without token → 401.

10.4.2 [S4-F11] Record Contract Milestone Event

Branch: feat/contract/S4-F11/<ID>

Endpoint: POST /api/contracts/{id}/milestones/track

Auth: Required (USER)

Databases: Cassandra, PostgreSQL, MongoDB

Request body:

```
{
  "milestoneOrder": 2,
  "status": "COMPLETED",
  "recordedBy": "Ahmed Ali",
  "notes": "Design phase deliverables approved by client"
}
```

Behavior:

- a) Validate the JWT token –**throws 401 if missing or invalid.**
- b) Find contract by ID in PostgreSQL –**throws 404 if contract not found.**
- c) Validate status is one of PENDING, IN_PROGRESS, COMPLETED, APPROVED (matches the M1 ProposalMilestone.status enum) –**throws 400 if the status value is not one of those.**
- d) Store a new milestone event in the Cassandra time-series table with the contract ID as the partition key, the current timestamp as the clustering column, and the milestone order, status, recordedBy name, and notes from the request. Cassandra is the **primary storage** for the milestone time-series.
- e) **Additionally**, fire the Observer chain to log a MILESTONE_TRACKED event to the contract_events MongoDB collection with the contractId, the milestoneOrder, the status, and a summary of the payload in details. Both writes must succeed independently – Cassandra for the time-series trail, MongoDB as the observability audit log that feeds the activity feed.
- f) Return status **code 201** indicating the event was recorded.

Test scenario:

- a) Create a contract via CRUD. POST /api/contracts/{id}/milestones/track with status=PENDING, milestoneOrder=1 → 201. Verify a row exists in Cassandra contract_milestone_events for this contractId, AND a MILESTONE_TRACKED document exists in MongoDB contract_events.
- b) Record another event with status=IN_PROGRESS, milestoneOrder=1 → 201. Query the timeline (S4-F12) → should show both events in reverse chronological order.
- c) POST /api/contracts/999/milestones/track → 404.
- d) Call with status="BANANA" → 400.
- e) Without token → 401.

10.4.3 [S4-F12] Get Contract Milestone Timeline

Branch: feat/contract/S4-F12/<ID>

Endpoint: GET /api/contracts/{id}/milestones/timeline?startTime={datetime}&endTime={datetime}

Auth: Required (USER)

Databases: Cassandra, PostgreSQL, Redis

Response: List of ContractMilestoneDTO: timestamp, milestoneOrder, status, recordedBy, notes.

Behavior:

- a) Validate the JWT token –**throws 401 if missing or invalid.**
- b) Find contract by ID in PostgreSQL –**throws 404 if contract not found.**
- c) Query the Cassandra time-series table for all milestone events for this contract, optionally filtered by the startTime and endTime parameters. If no time range is provided, return all events. Results are sorted by most recent first (Cassandra clustering order).
- d) The response should be cached for 5 minutes.
- e) Return the milestone timeline with status **code 200**. Return an empty list if no events exist.

Test scenario:

- a) Create a contract. Record 3 milestone events: PENDING at 14:00, IN_PROGRESS at 14:15, COMPLETED at 14:30. GET /api/contracts/{id}/milestones/timeline → should return 3 events with COMPLETED first.

- b) Query with `startTime=14:10&endTime=14:20` → should return only the `IN_PROGRESS` event.
- c) `GET /api/contracts/999/milestones/timeline` → 404.
- d) Contract with no milestone events → empty list.
- e) Without token → 401.

10.5 Wallet Service Features (port 8085)

10.5.1 [S5-F10] Get Platform Fee Analytics by Job Category

Branch: `feat/wallet/S5-F10/<ID>`

Endpoint: `GET /api/payouts/analytics/category?startDate={date}&endDate={date}`

Auth: Required (USER)

Databases: PostgreSQL, MongoDB, Redis

Response: List of `CategoryRevenueDTO`: `category`, `netPayoutRevenue`, `platformFeeRevenue`, `totalRevenue`, `payoutCount`.

Platform fee source: S5-F10 reads `transactionDetails.platformFee` written by M1's updated S5-F4. See Section 4.6 for the full specification of this additive JSONB key and the 10%-of-amount fallback rule for pre-M2 Payouts that do not carry the key.

Behavior:

- a) Validate the JWT token –**throws 401 if missing or invalid**.
- b) Validate date range –**throws 400 if startDate is after endDate**. Both `startDate` and `endDate` are `LocalDate` values on the wire; internally expand the filter window to the fully-closed range [`startDate T00:00:00`, `endDate T23:59:59.999`] in the server time zone, matching the PG `TIMESTAMP` columns (`createdAt`, etc.) and the MongoDB `timestamp` field. Boundary rows (events exactly at `startDate T00:00:00` or `endDate T23:59:59.999`) are included.
- c) Use the M1 cross-service native SQL pattern: JOIN payouts with contracts (on `contracts.id = payouts.contract_id`) and jobs (on `jobs.id = contracts.job_id`) on the shared PostgreSQL database. **Do not introduce inter-service HTTP calls**; those come in Milestone 3.
- d) Filter by `payouts.createdAt` within [`startDate`, `endDate`] and `payouts.status = COMPLETED`.
- e) Group by `jobs.category`. For each group compute:
 - `platformFeeRevenue` = sum of `transactionDetails->'platformFee'` cast to numeric (default to 10% of the payout amount when the key is missing).
 - `netPayoutRevenue` = sum of `payouts.amount` – `platformFeeRevenue`.
 - `totalRevenue` = sum of `payouts.amount`.
 - `payoutCount` = count of distinct payout IDs.
- f) Log an `ANALYTICS_VIEWED` event to the `payout_audit_trail` collection in MongoDB. **This log must be written on every invocation**, independently of whether the response was served from cache — perform the logging step **outside** the cache decorator/layer so it runs on cache hits too. Note that `ANALYTICS_VIEWED` / `DASHBOARD_VIEWED` are pure-observability actions and are excluded from Observer-driven cache invalidation (see §4.4.4) — they do not mutate the data that this dashboard reads, so invalidating the cache on them would defeat the cache.
- g) The response should be cached for 10 minutes.
- h) Return the breakdown with status **code 200**. Return an empty list if no data exists for the date range.

Test scenario:

- a) Create jobs: J1 (WEB_DEV), J2 (MOBILE), J3 (WEB_DEV). Create contracts and payouts: 3 payouts on WEB_DEV jobs totaling 6000 (with platform fees 600), 2 payouts on MOBILE jobs totaling 4000 (with platform fees 400). GET analytics/category → WEB_DEV: netPayoutRevenue=5400, platformFee=600, total=6000, count=3. MOBILE: netPayoutRevenue=3600, platformFee=400, total=4000, count=2.
- b) Date range with no payouts → empty list.
- c) Invalid date range → 400.
- d) Without token → 401.

10.5.2 [S5-F11] Get Payout Method Breakdown

Branch: feat/wallet/S5-F11/<ID>

Endpoint: GET /api/payouts/analytics/methods?startDate={date}&endDate={date}

Auth: Required (USER)

Databases: MongoDB, PostgreSQL, Redis

Response: List of PayoutMethodDTO: method, successCount, failureCount, successRate, totalAmount.

Behavior:

- a) Validate the JWT token –**throws 401 if missing or invalid**.
- b) Validate date range –**throws 400 if startDate is after endDate**. Both startDate and endDate are LocalDate values on the wire; internally expand the filter window to the fully-closed range [startDate T00:00:00, endDate T23:59:59.999] in the server time zone, matching the PG TIMESTAMP columns (createdAt, etc.) and the MongoDB timestamp field. Boundary rows (events exactly at startDate T00:00:00 or endDate T23:59:59.999) are included.
- c) Query the payout_audit_trail collection in MongoDB filtered by timestamp within the fully-closed range [startDate, endDate] (inclusive on both ends; interpret both as naive LocalDateTime, same server time zone as the PG created_at columns). Consider only events with action ∈ {COMPLETED, FAILED} — exclude CREATED, REFUNDED, REFUND_DENIED, and ANALYTICS_VIEWED from the method-breakdown counts.
- d) Group the filtered events by method field (BANK_TRANSFER, PAYPAL, CRYPTO, matching the M1 Payout.method enum). For each method:
 - successCount = count of events where action == COMPLETED
 - failureCount = count of events where action == FAILED
 - successRate = successCount / (successCount + failureCount) (return 0 if denominator is 0)
 - totalAmount = sum of amount across the COMPLETED events only (do not sum failed amounts)
- e) The response should be cached for 10 minutes.
- f) Return the breakdown with status **code 200**. Return an empty list if no data.

Test scenario:

- a) Create payout audit events: 5 successful BANK_TRANSFER payouts totaling 5000, 2 failed BANK_TRANSFER payouts, 3 successful PAYPAL payouts totaling 3000. GET analytics/methods → BANK_TRANSFER: success=5, failure=2, rate≈0.71, total=5000. PAYPAL: success=3, failure=0, rate=1.0, total=3000.
- b) Date range with no events → empty list.
- c) Invalid date range → 400.
- d) Without token → 401.

10.5.3 [S5-F12] Process Milestone-Based Payout Reversal

Branch: feat/wallet/S5-F12/<ID>

Endpoint: POST /api/payouts/{id}/reverse-milestone

Auth: Required (USER)

Databases: PostgreSQL, MongoDB, Redis

Design Pattern: Strategy (see Section 3.2)

Request body:

```
{
  "reason": "milestone_dispute",
  "reversalScope": "MILESTONE_ONLY"
}
```

Note on path: This is a **distinct endpoint** from M1's S5-F2 (PUT /api/payouts/{id}/refund). M1's endpoint performs a simple full refund and remains unchanged; M2's new /reverse-milestone endpoint selects a reversal strategy at runtime (see Section 3.2) and writes a richer MongoDB audit trail. Both must coexist.

Behavior:

- a) Validate the JWT token –**throws 401 if missing or invalid**.
- b) Find payout by ID in PostgreSQL –**throws 404 if payout not found**.
- c) Validate the payout status is COMPLETED –**throws 400 if the payout is not in COMPLETED status** (you cannot reverse a payout that has not been completed).
- d) Select the reversal strategy (see Section 3.2 for the three required strategies) **without throwing yet**: FullPayoutReversalStrategy (when reversalScope=FULL and the payout is within the 30-day window from createdAt), MilestoneReversalStrategy (when reversalScope=MILESTONE_ONLY and within the window), or NoReversalStrategy (when the payout is older than 30 days from createdAt).
- e) **If the selected strategy is NoReversalStrategy**: (i) Log a REFUND_DENIED event to the payout_audit_trail collection in MongoDB with the strategy name NoReversalStrategy and the denial reason “reversal window expired”. (ii) Invalidate wallet-service::S5-F10::*, wallet-service::S5-F11::*, and wallet-service::payout:::{id} (per Â§4.4.4 NoSQL-writer + PG-writer rules) – even though no PostgreSQL row mutates on the denial path, the new REFUND_DENIED audit event changes what those analytics would return on the next read. Over-invalidation is acceptable per Â§8.2. (iii) Then **throw 400 with message “reversal window expired”**. Logging and cache invalidation must both happen **before** the throw so the audit event is persisted and cached reads pick up the new event immediately.
- f) Otherwise (strategy is FullPayoutReversalStrategy or MilestoneReversalStrategy): invoke the strategy's calculateRefund(payout, request) to obtain the reversal amount. For MilestoneReversalStrategy, the amount is the sum of ProposalMilestone.amount values whose status is NOT in {COMPLETED, APPROVED} (resolved via the cross-service native SQL pattern: contract → proposalId → proposal_milestones). Update the payout status to REFUNDED.
- g) Record the reversal in the transactionDetails JSONB column of the Payout (reusing the M1 JSONB pattern from S5-F2). Add keys: refundAmount, reversalScope, refundReason, and refundedAt. Do **not** add a new SQL column.
- h) Log a REFUNDED event to the payout_audit_trail collection in MongoDB with the strategy name, reason, original amount, reversal amount, and reversal scope.
- i) Invalidate wallet-service::S5-F10::*, wallet-service::S5-F11::*, and wallet-service::payout:::{id} (per Â§4.4.4 NoSQL-writer + PG-writer rules). Over-invalidation is acceptable per Â§8.2.

- j) Return the updated payout with status **code 200**.

Test scenario:

- a) Create a completed payout of 5000 whose contract's proposal has 3 milestones (2 COMPLETED at 1500 and 1500, 1 IN_PROGRESS at 2000), created today. POST /api/payouts/{id}/reverse-milestone with reversalScope: "MILESTONE_ONLY" → 200, refund amount = 2000 (only the incomplete milestone), status = REFUNDED. Verify transactionDetails.refundAmount=2000 and reversalScope="MILESTONE_ONLY".
- b) Create another completed payout of 3000, created today. Reverse with reversalScope: "FULL" → 200, refund amount = 3000. Verify transactionDetails.refundAmount=3000 and reversalScope="FULL".
- c) Create a payout with createdAt 40 days ago → 400 with message "reversal window expired" (NoReversalStrategy selected). Verify that a REFUND_DENIED document exists in payout_audit_trail for this payoutId (audit log is written **before** the 400 is thrown). Also verify that any previously-cached key under wallet-service::S5-F10::* and wallet-service::S5-F11::* was removed from Redis (cache invalidation on the denial path).
- d) Attempt to reverse a PENDING payout → 400.
- e) Attempt to reverse an already REFUNDED payout → 400.
- f) POST /api/payouts/999/reverse-milestone → 404.
- g) Without token → 401.